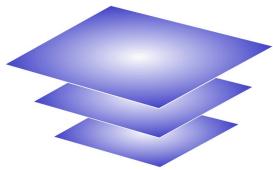


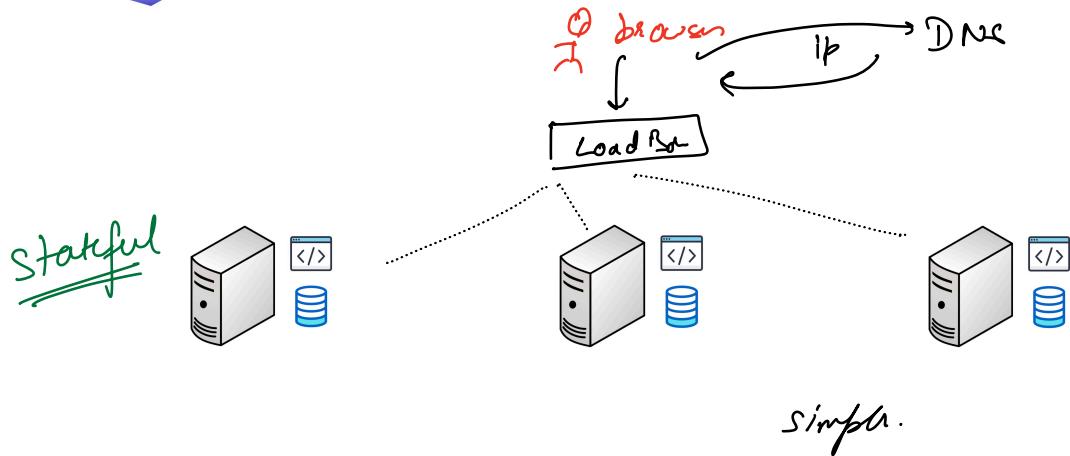
Agenda



1. Appserver Layer support@scaler.com
2. Caching
3. Invalidation & Eviction 10:30
4. Local Caching



Appserver Layer



?

Is this good design?

1. Data & business logic are tightly coupled.
2. for any code deployment, data temporarily unavailable.
3. webserver
 - ↳ CPU
 - ↳ RAM
 - ↳ fast I/Odatabase
 - ↳ disk space
 - ↳ fast I/O

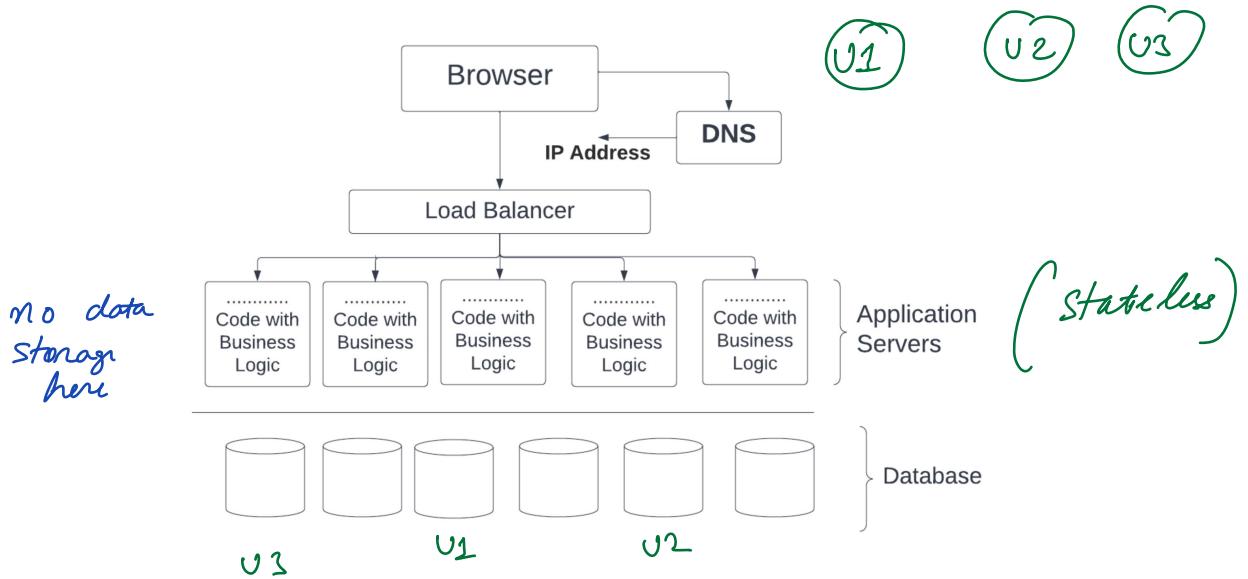
Split it!

Server is expensive
webserver & db are competing for resources.

?

Cons

1. Complex (buy more albeit specialized hardware)
 - WEB
2. additional latency.
 - DB
 - DB
 - DB off
 - DB
 - DB
 - DB other...



? Which request to which server? No. ∵ opp servers are stateless

exceptions → authentication & authorization
(session)
(causes when you need stateful)
go True

→ maintain local cache



Who does the Consistent Hashing? Load Balancer / App Server

X stateless app servers	✓ ∵ database is stateful & sharded.
-------------------------------	---

Java Client
MySQL client

→ implement consistent hashing
informally.



How do the Appservers know how many & which DB servers are available?

Zookeeper

cluster manager



Making Tea

? Ingredients

tea leaves / powder

water milk

sugar

Earl grey tea



supermarket



fridge

Caching layer-

→ slow access

→ access is very fast

→ source of ingredients
v.v. wide selection

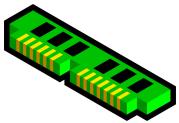
→ limited in space
only have freq. used items.

→ always has fresh items available

→ items can go stale

multiple levels of cache





Caching



→ Cookies / local storage / indexedDB

User name
theme (light/dark)
config
profile pic
language



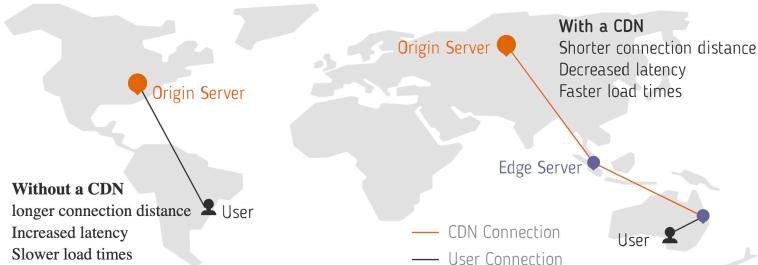
- ① user data
- ② images/videos / static media
- ③ css/js/html
- ④ IP address



googh → USA

Cloudflare 375,000 servers
Akamai
Amazon Cloudfront
Fastly

Content delivery networks



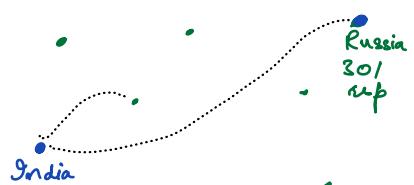
<https://www.cloudflare.com/learning/cdn/what-is-a-cdn/>

? How to get the IP of the nearest CDN server?

↳ geo DNS (less adopted)

↳ anycast

DNS maps
cdn.cloudflare.com



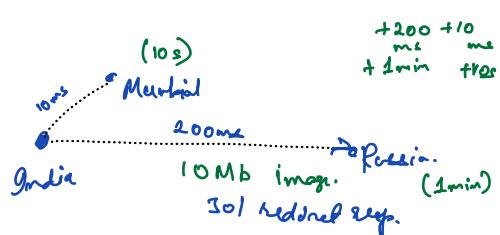
<https://www.cloudflare.com/en-gb/learning/cdn/glossary/anycast-network/>

<https://constellix.com/news/anycast-vs-geodns>

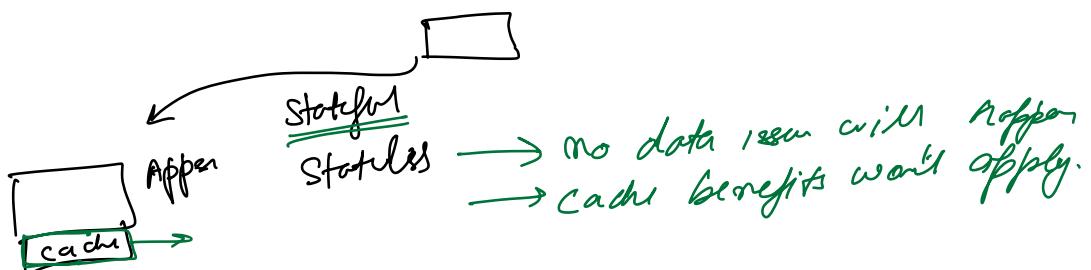
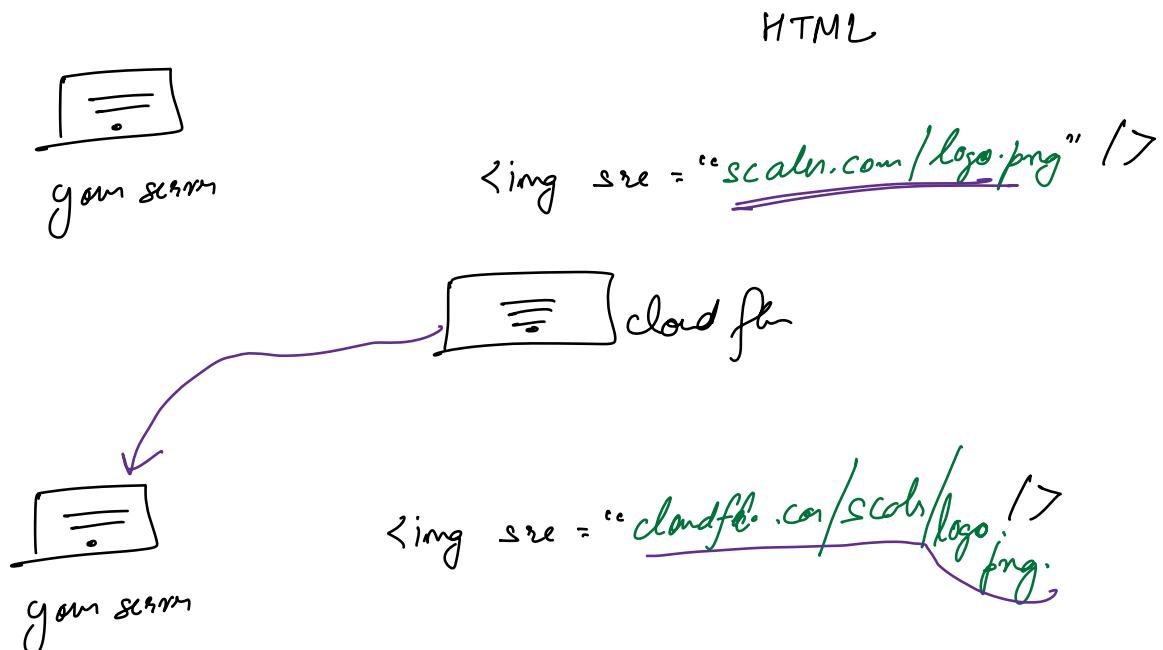
Local Cache each appserver maintains its own cache

Global Cache cache layer shared by all appservers.

→ url → ip



How do CDNs get data?





Challenges with Caching



- Cache is not the source of truth (data)
- data in cache could be stale
- periodically invalidate stale data



- limited in size
- evict data from cache



Cache Invalidation



Write Through

Write Back

Write Around



Cache Eviction

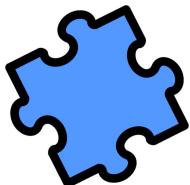
First in First Out

Least Recently Used

→ Hashmap + LRU
↳ Hashmap + DLL

Last in First Out

Least Frequently Used



For next class



Local caching of test data on the app server. How to invalidate on the update?



How is facebook's newsfeed calculated, and how can you make it very fast?



Some hosting where media files are stored
HTML CSS JS

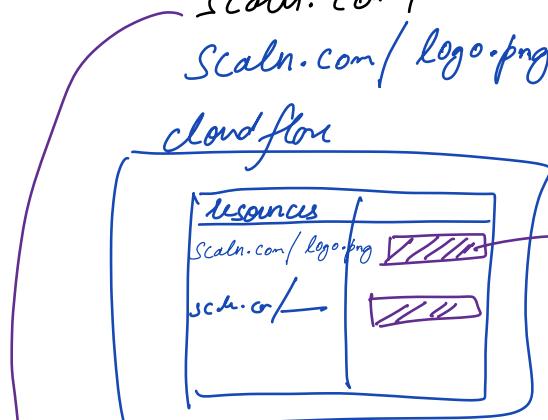
scdn.com/index.html

Scdn.com/logo.png

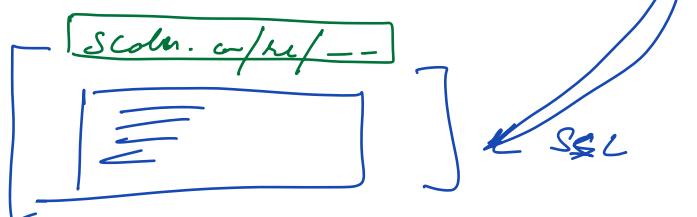
→ vercel

image video

→ vercel



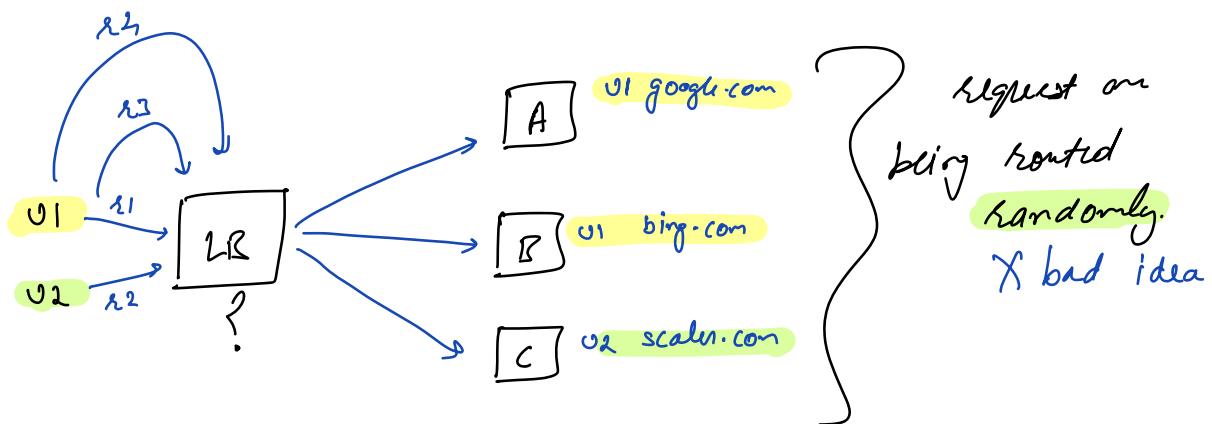
cdn.cloudflare.com/1/abe57



① multiple app-servers
 (assume data is stored in app-servers themselves)

② data is sharded
 (partitioned across multiple servers)

③ LB gets requests from user & needs to send them to an appserver.



if user wants to read data LB has no idea!!

$(\text{User-id \% } N)$

Range band

the user's req go to same server
 LB knows which server to send to.
 # servers increases or decreases.
 ↳ lot of data shuffling.

hash functions

$$H_1(x) = (10x^2 - 3x + 1) \% 100$$

$$H_2(x) = (x^3 + 3x^2 - 2x) \% 100$$

$$H_3(x) = (3x^2) \% 100$$

H_1

v_1 77

v_2 38

v_3 12

v_4 56

v_5 22

	H_1	H_2	H_3
s_1	37	2	10
s_2	95	99	73
s_3	0	11	58

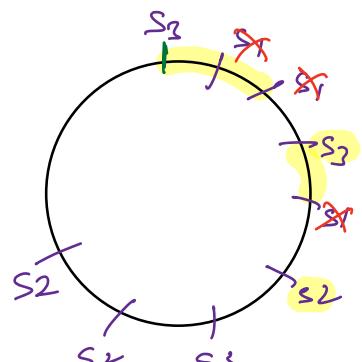
Load Balancer

- ① Health check / heartbeat to find out what all servers are alive.

$s_1 \ s_2 \ s_3$

②

0	2	10	11	37	45	58	73	99
s_3	s_1	s_1	s_3	s_1	s_2	s_3	s_2	s_2



- ③ v_1 req.

$$H_1(v_1) = 77 \rightarrow s_2$$

v_2 req.

$$H_1(v_2) \rightarrow s_2$$

v_3 has

$$H_1(v_3) = 12 \rightarrow s_1$$

- ④ add new server s_4
 Load Balancer detects that s_3 is also full.

$$\begin{array}{ccc} H_1 & H_2 & H_3 \\ s_1 & 76 & 16 \\ s_2 & & 55 \end{array}$$

Reconstruct.				s1				s2			
0	2	10	11	16	37	45	58	73	99		
s_3	s_1	s_1	s_3	s_1	s_2	s_3	s_3	s_2	s_2	s_4	

H_1

$$v_1 \quad 77 \rightarrow s_2$$

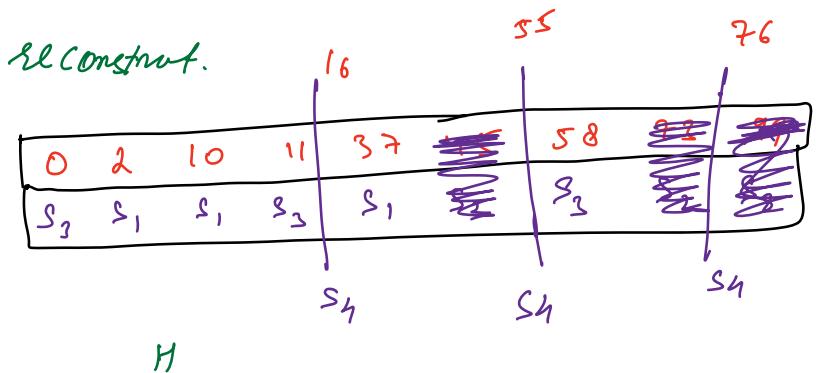
$$v_2 \quad 38 \rightarrow s_2$$

$$v_3 \quad 12 \rightarrow \text{earlier } s_1 \text{ now } s_4$$

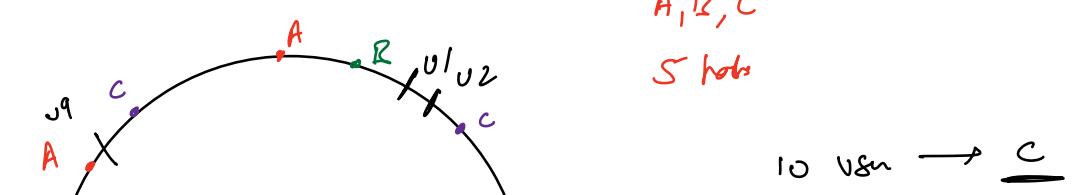
$$v_4 \quad 56 \rightarrow s_2$$

$$v_5 \quad 22 \rightarrow s_1$$

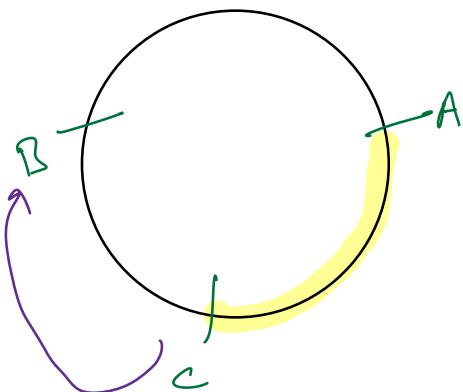
- ⑤ s_2 goes down
 load balancer detects that s_2 is down



v_1 77 $\rightarrow S_2$
 v_2 38 $\rightarrow S_2$
 v_3 12 \rightarrow earliest S_1 move S_3
 v_4 56 $\rightarrow S_2$
 v_5 22 $\rightarrow S_1$



$v_1 \rightarrow A$
 $v_2 \rightarrow A$
 $v_3 \rightarrow B$
 $v_4 \rightarrow B$
 $v_5 \rightarrow B$
 $v_6 \rightarrow B$
 $v_7 \rightarrow B$
 $v_8 \rightarrow A$
 $v_9 \rightarrow A$
 $v_{10} \rightarrow A$



with 1 hash
load of $S_2 \rightarrow$ just
 X one
serv.

100 servers each serve $1\% \text{ of total load}$
 S_{43} goes down $\rightarrow S_{57}$
 eating (1% norm) cascading failure
 $1 - \frac{43}{57} = \frac{14}{57} \approx 25\%$
 $25\% + 1\% = 26\%$

S_{43} 16 hashes $\Rightarrow 16 \text{ diff serv}$
 $99 - 96 \Rightarrow 1\%$
 $16 \Rightarrow 1\% + \frac{1}{16}$
 1.0625%

