

Agenda



1. Invalidation & Eviction
2. Local Caching - Case study 1
3. Global Caching - Case study 2
4. Case study 3 - Facebook News Feed

support@scaler.com

9 - 11:30 class

+10 mins

doubt

Most applications are read heavy → caching.



Challenges with Caching



- Cache is not the source of truth
↳ database
- Cache data can get stale

- Small in size
- Cache miss → data wasn't found in cache
↳ lead to slowdown
- to load new data
Cache needs to remove existing data.
↳ what data to remove?



Cache Invalidation

↳ fix stale data

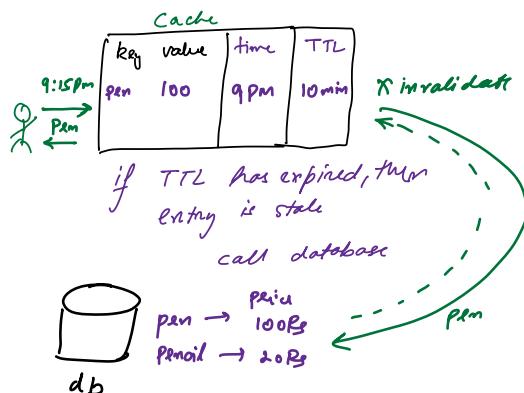


Time to Live (TTL)

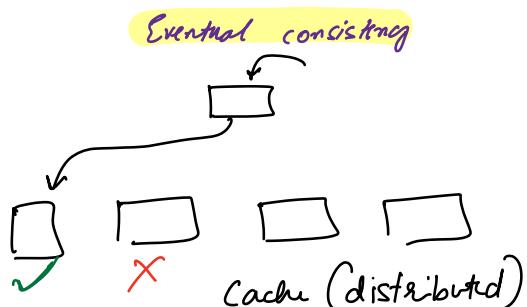
Simple / Configurable TTL

you will get slightly stale data

Write Through



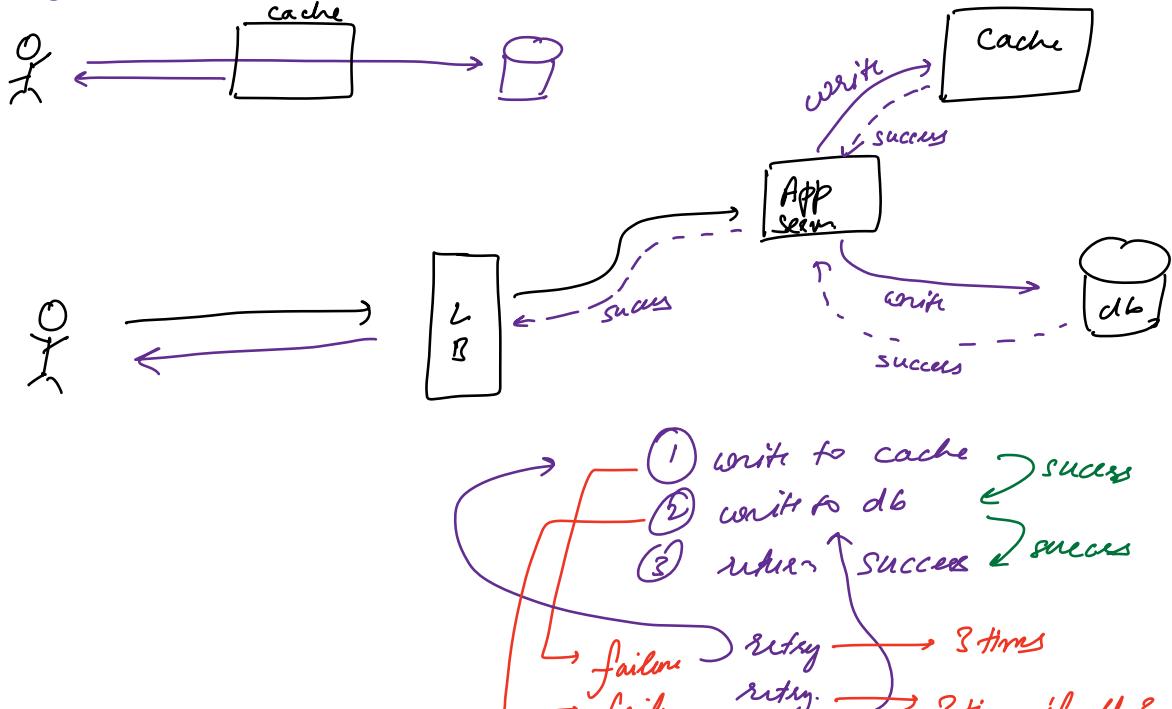
Write Back



Write Around

Write Through Cache

Any update must first be written to cache.



? Can cache data be stale
in case of Write Through Cache?

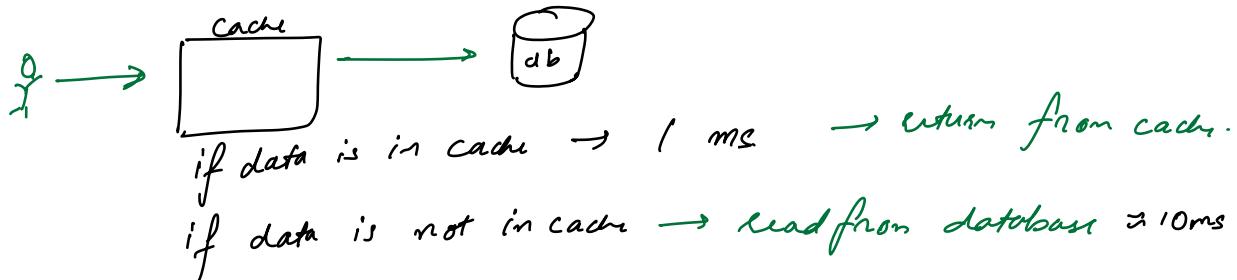
Never!!

- ① Always fresh data
- ② Reads are super fast → every cache hit will return directly from cache
- ③ writes are much slower

Write through cache is good for a **read-heavy** system.

most req. are reads & not writes
90% of applications

Reads are fast when using a cache
↳ ∵ of how cache works.



Hit → 1 ms → 99% cache hit

miss → 1 ms + 10 ms = 11 ms → 1% misses

$$\begin{aligned} \text{avg. time per query.} &= 0.99 * (1 \text{ ms}) + 0.01 * (11 \text{ ms}) \\ &= 0.99 \text{ ms} + 0.11 \text{ ms} \\ &= 1.1 \text{ ms} \end{aligned}$$

Write around cache

→ write happens directly at the database

↳ the cache might be out of sync for a short while

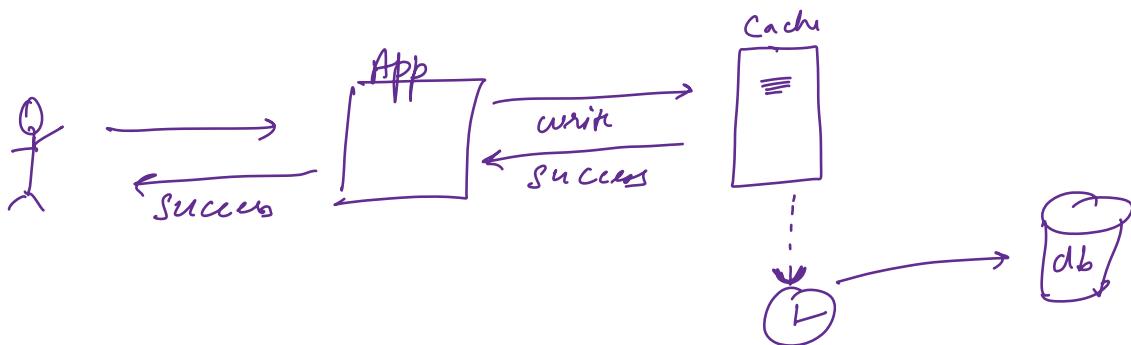
→ cron job that queries recently updated data from db (10 mins) & updates in cache.

select * from product-price
where updated_at > (now() - 10m)
} → update in cache.

Write-back cache (inconsistent)

↳ first write data in cache
↳ return success
even though data is not on db!!

↳ cron job that will periodically flush
this data on to the database.



Cons

if cache server crashes we lose some data!!

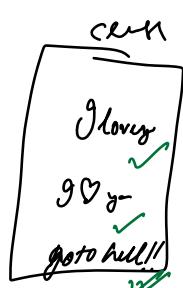
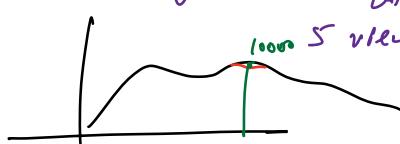
Pros

Extremely high throughput

Use write-back cache in write-heavy system
only if consistency is not important

Analytics / tracking clicks or ads

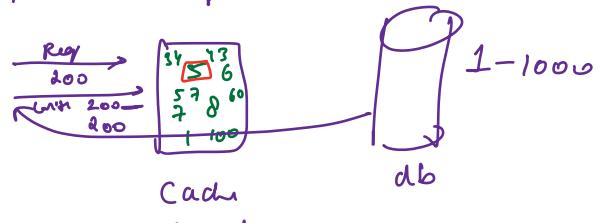
↓
trend based.
↳ 10,000 views
↳ only 9,995 views got credit
↳ 5 view → lost





Cache Eviction

Cache is limited in size



First in First Out

Simple
Simpler



Least Recently Used

time
Hashmap + Heap
Hashmap + DLL

if you see something, it is v. likely
that you will see it again in
near future.

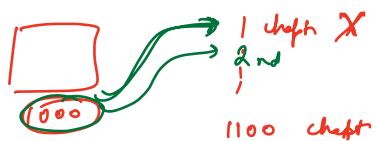
Last in First Out

↳ whatever you've read recently will get removed.
HW → find out circumstances when this is awful.

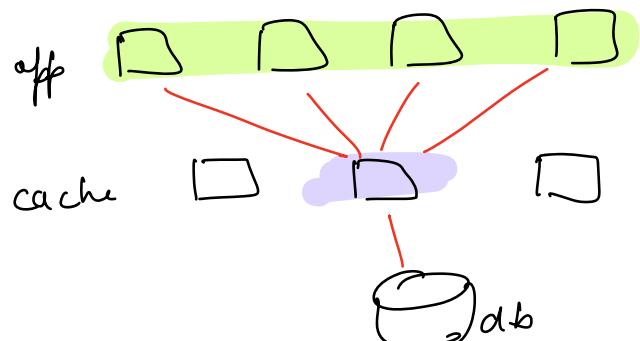
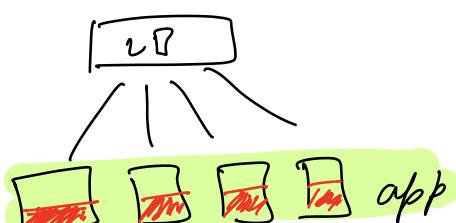
Least Frequently Used

evicted.
read counts

One Piece Meme: Reading groups.



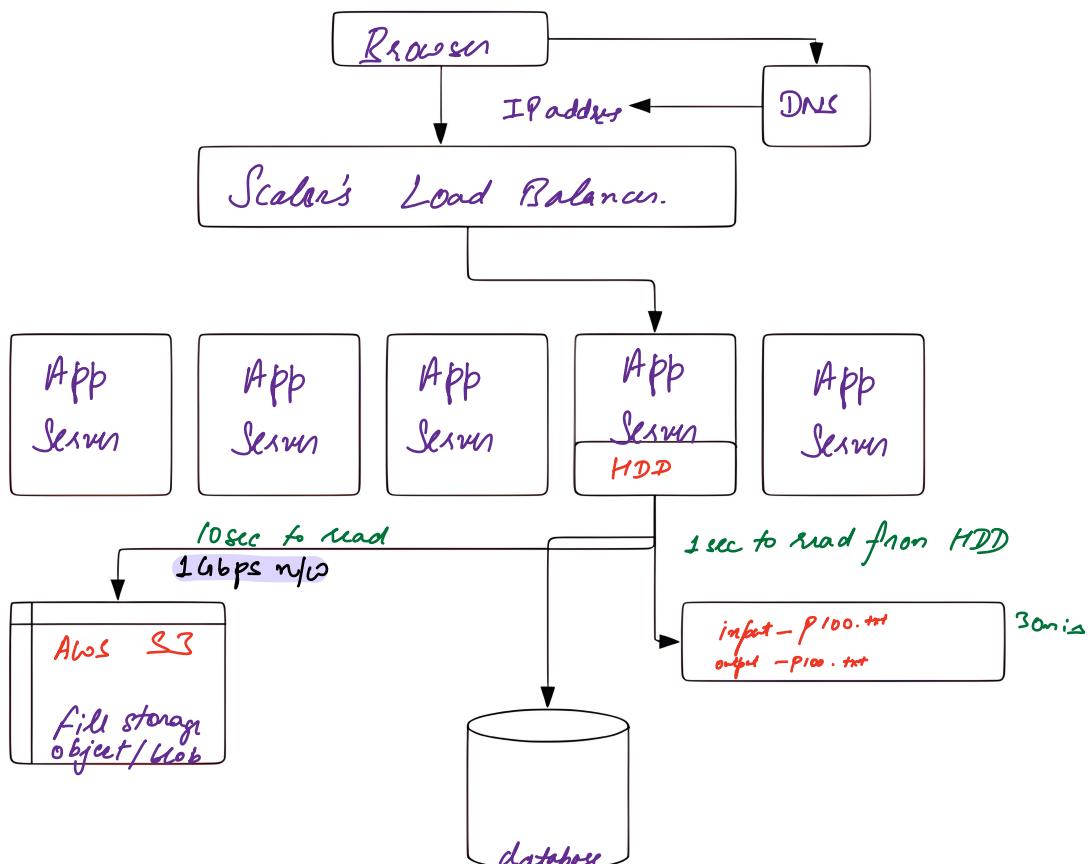
used by only
1 appserver.



CASE STUDY



Online Code Judge



<https://gist.github.com/jboner/2841832>

? What does the AppServer need to evaluate a solution?

problem-id problem-metadata

user-id solved-already score

code language. **input testcases & output testcases**

? How large can this data be?

input
output } 500 MB each → 1 GB of testdata
for a single problem!

given array → sort it
 $O(n \lg n)$

$\approx 10^6$ size

40-50 testcases

500MB input.
500MB output



Change to the testcases

if a test case is changed → how do we update it in the local cache?

- at every write, push change to all app servers.

↳ write through cache



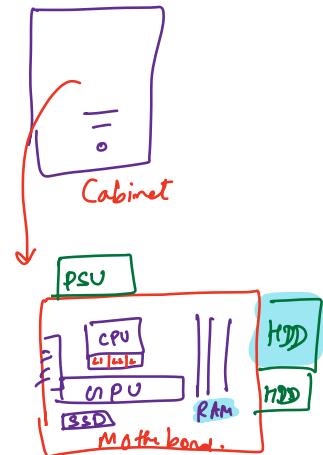
→ every app server has its own cache.

→ update every app server

→ wasteful

↳ write back cache & TTL

some period during which data will be stale.



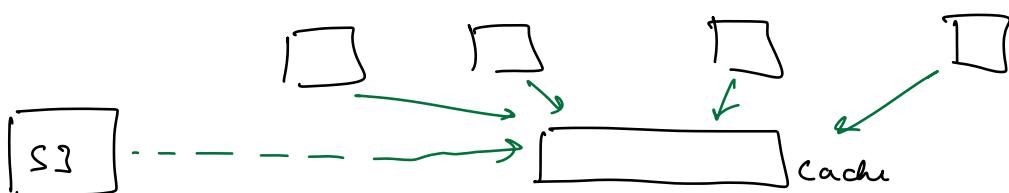
- TTL / write back cache.

bad solution → we want immediate effect

10mins TTL → entire content gets over & test case are not updated

30sec TTL → most req will have to go to S3
no benefit of cache

- Global Cache

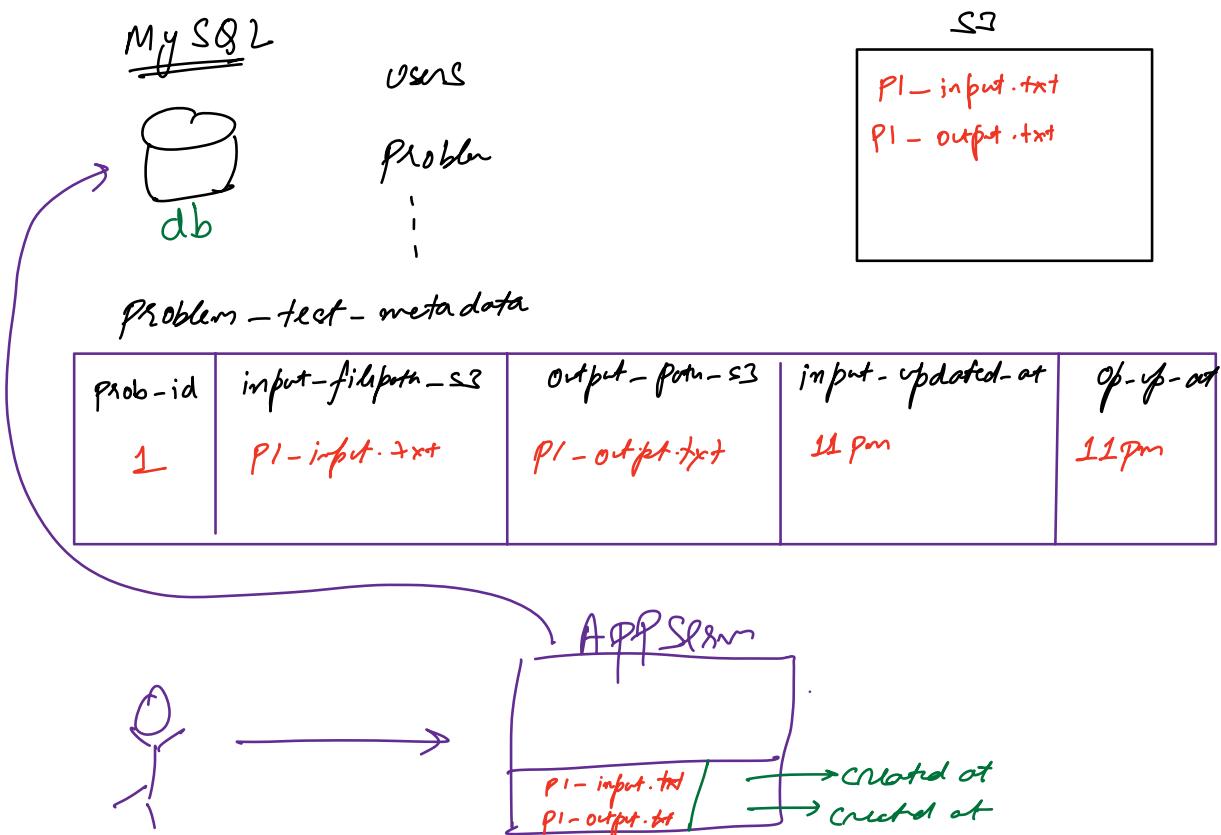


- ① a lot of m/w traffic will come to cache
- ② if cache seem has 32GB RAM
↳ 32 files

a lot of cache misses.

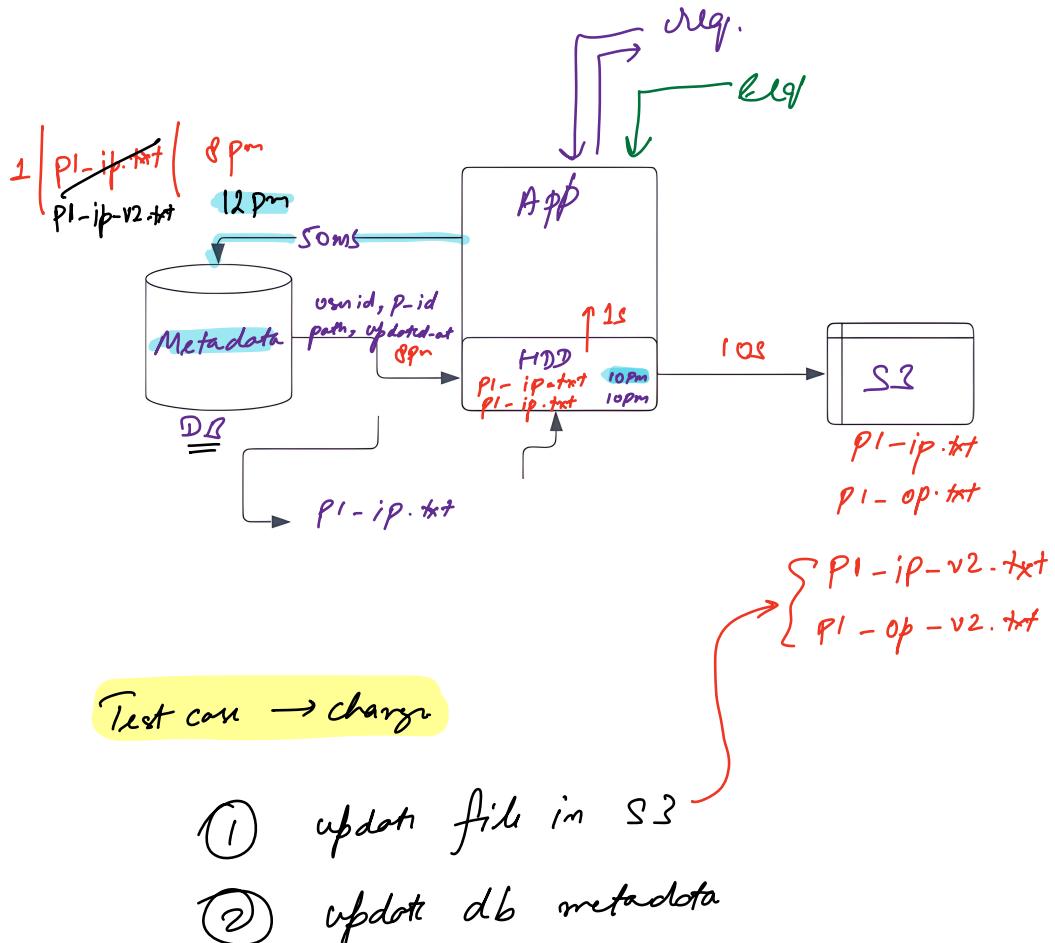
- ③ store in HDD → SSD ↳ $\sim 300 \text{ MBps}$
- HDD

File Metadata in DB





Solution





Contest Leaderboard

? How can we compute the ranklist?

? How frequently does the ranklist change?

? Is immediate consistency important?



Local cache



Global cache



try.redis.io
redis.io/docs/data-types
university.redis.com



Facebook News Feed

? How can we compute the News Feed?

Profile Page

News Feed

?

Does all the user data fit on a single machine?

?

What is my sharding key?

?

How do I retrieve data for the news feed?

?

How can you optimize the construction of the News Feed?



Cache user -> news feed

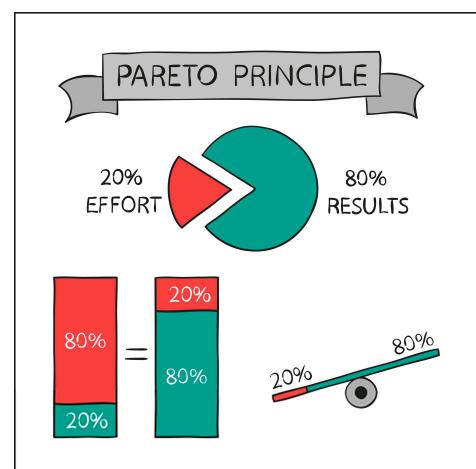
1.

2.

3.



Back of the envelope estimates

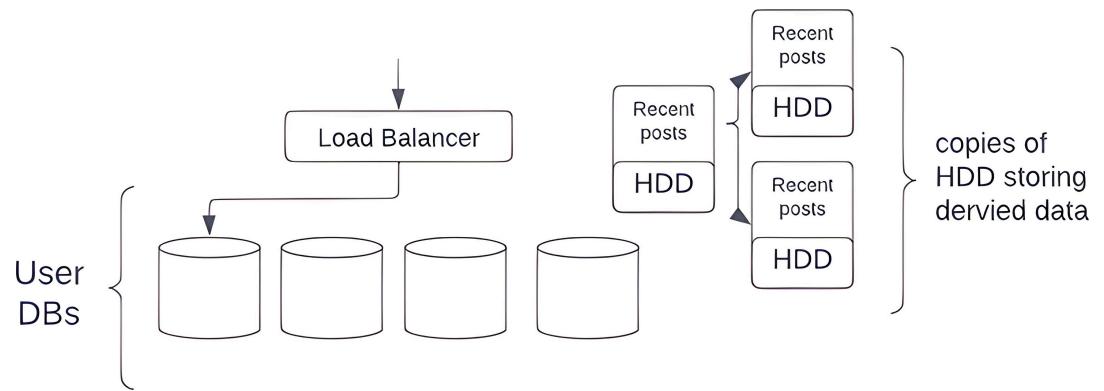


?

Daily post data

?

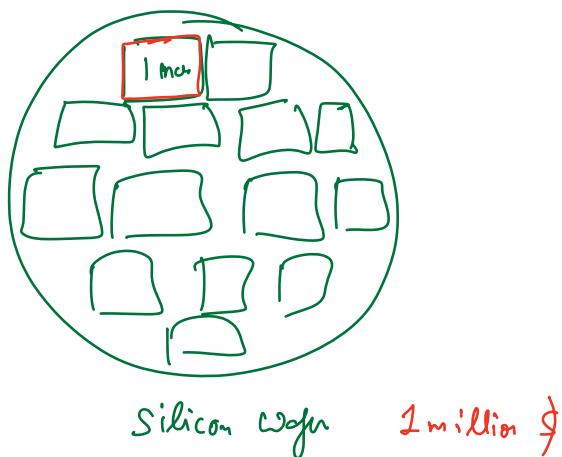
Monthly post data





CPU → Intel
clean rooms

ULS I
Ultra large scale integration
3nm
5nm
V.V.V. expensive



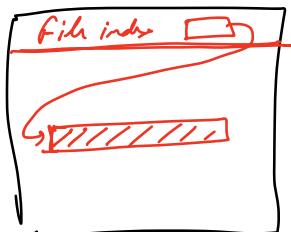
flip-flop → registers 50Mb
→ cheap but slow L1
→ dear but slow L2
;

- ① write through → no stale data
- ② write around → cron job that updates cache using db
 \approx TTL
stall for some time
- ③ write back → inconsistent data

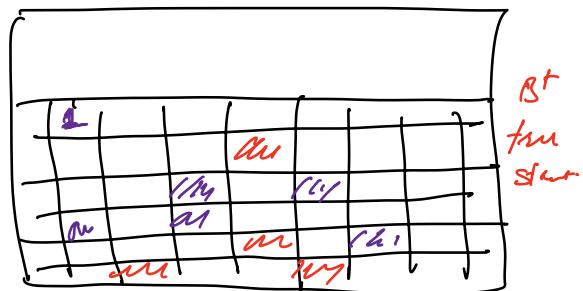
meta data → db metadata → same if true? did we have stale data
No stale data!



S3
↓
NTFS / CxTz
fill storage

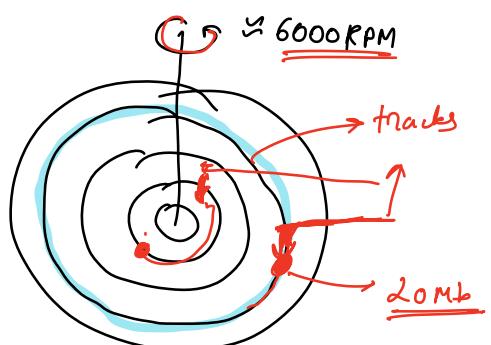
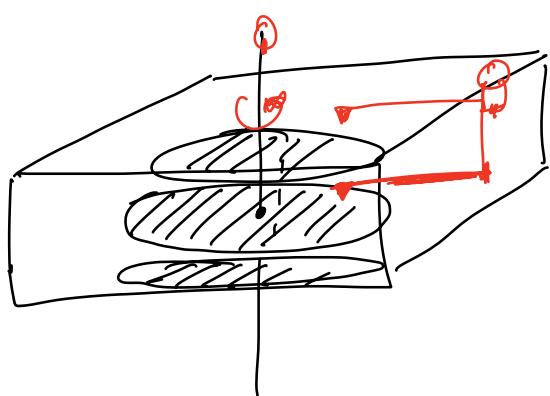


DB querying data



continuous data access

P-id | input soon / out soon
random data access



$$\text{disk seek} \Rightarrow 100\text{ms} \quad \frac{6000}{60\text{s}} = 10\text{ms}$$

$TTL \rightarrow$ stale data

global \rightarrow high m/o usage

\hookrightarrow HDD \rightarrow slow

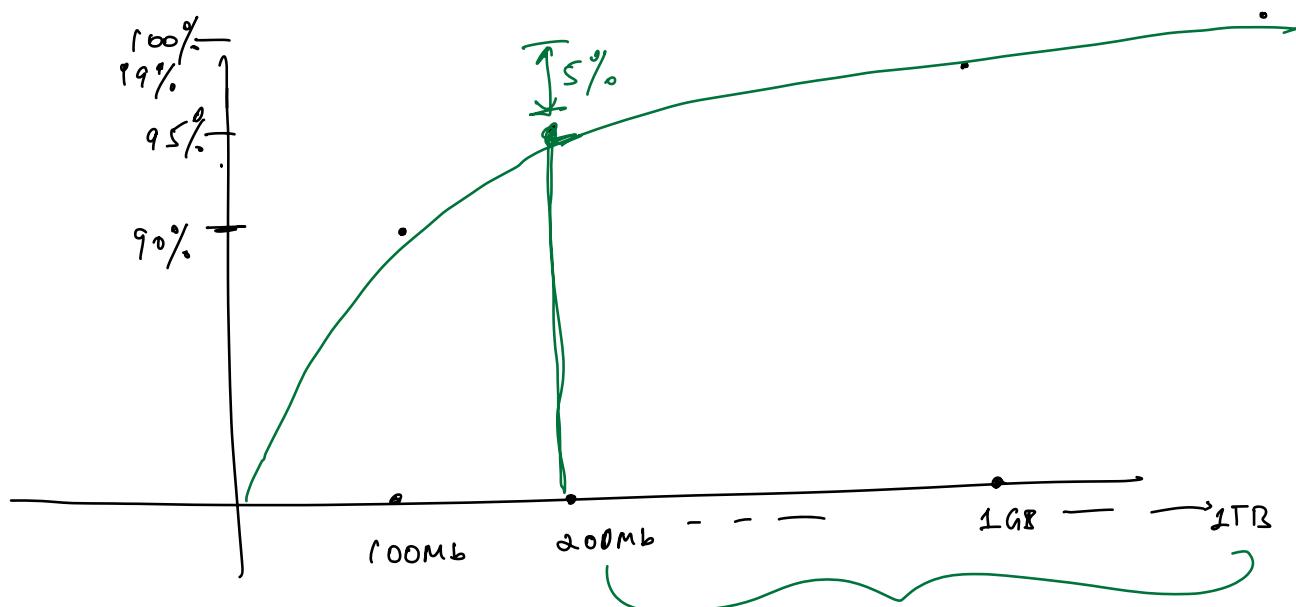
\hookrightarrow RAM \rightarrow small

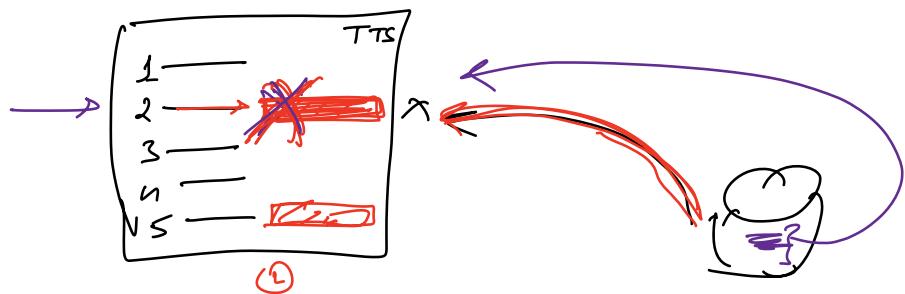
written back $\times \rightarrow$ inconsistent

written around \rightarrow similar to TTL

DB \rightarrow 1 TB

Cache \rightarrow 100 Mb





- ① wait for req (2) → expired → invalidate
 - ② Redis detected that (2) is still → remove it
 - ③ cron → stale data → refetch from db to update in cache
-

