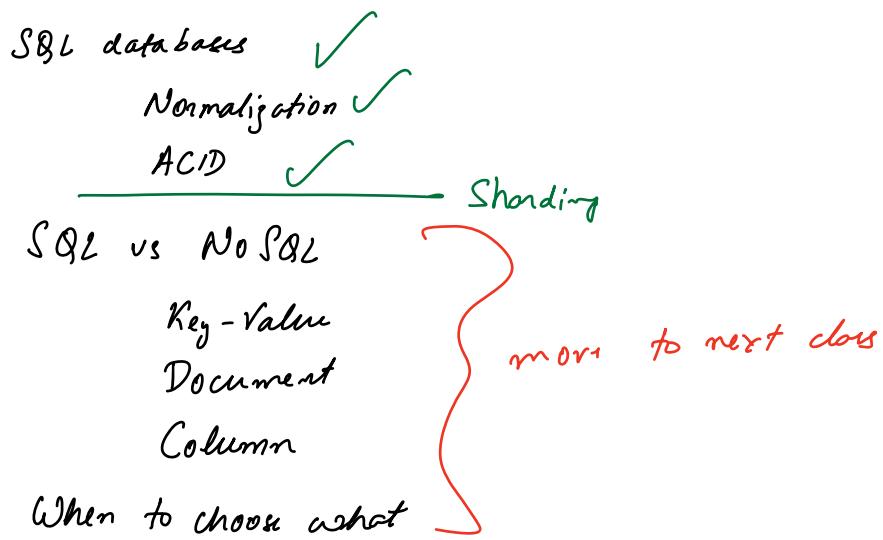


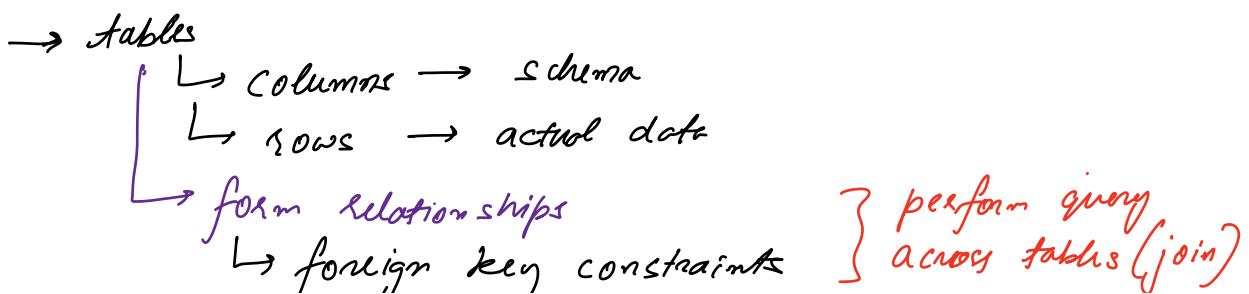
## Agenda



NoSQL  $\Rightarrow$  Say NO to SQL  $\rightarrow$  bad X  
Not Structured Query Language X  
Not only SQL ✓  
↳ if SQL does not satisfy your needs  
↳ data is just too large for SQL to handle

## The good parts of SQL

↳ Relational database



✓ ✓ ✓

→ **Normalization** → possible & easy

- ↳ removing redundancy
- ↳ achieves data consistency across tables

### Address

User-id	Pin-code	street-address	city	state
1	1000	—	Bangalore	Maharashtra
2	7000	—	Bangalore	Karnataka.
3	1009	—	Mumbai	Maharashtra
4	1009	Powai	Mumbai	Maharashtra.

### Duplicated data?

- ↳ wasted space
- ↳ anomalies (insert, update, delete)
  - ↳ either update all values in case of change
  - ↳ or deal with inconsistent data
  - ↳ deleting a row → unintended / side effects

### Scaln

student-id	student-name	course-id	course-name	course-length
1	Manu	5	HLD	1.5 months
2	Shaw	5	HLD	1.5 mn
3	Rajur	6	Data Engg.	2 months

Rajeev → only student enrolled in Data Engg.

normalize

normalization → process of de-duplicating data

<u>Student</u>		<u>Courses</u>		
id	name	id	name	duration
1	Maver	5	MLD	1.5 m
2	Shan	6	Data	2 m
X	Rajeev			

Student - courses / enrollments

FK student.id	FK course.id
student_id	course_id
1	5
2	5
3	6

→ ACID guarantees → out of the box in SQL  
↳ un-divisible

Atomicity → transactions are atomic  
either all or none → no in-between

Consistency → invariants are maintained before & after transaction

Isolation → transactions should NOT interfere  
no dirty reads

Durability → completed transactions are persisted  
even in case of failures (reboot/crash)

Rohit

~~1000Rs~~

100Rs

withdraw 900Rs → Neeraj

- fail ↗  
rollback ↘
- ① check if  $\text{rohit.balance} \geq 900$
  - ②  $\text{rohit.balance} - = 900$
  - ③  $\text{neeraj.balance} + = 900$
  - ④  $\text{neeraj.notify}()$

Before transaction

Rohit = 1000

Neeraj = 1000

Total = 2000

} magically  
vanished

After (failed) transaction

Rohit = 100

Neeraj = 1000

Total = 1100

in-consistent

Rohit → Neeraj

Audible ← charge 500Rs

1000  
100  
-900

Rohit  $\geq 900$  ✓  
Rohit -= 900  
Neeraj += 900  
Notify

9:41  
9:42  
9:43

dirty ← Rohit  $\geq 500$  ✓  
Rohit -= 500

9:41  
9:42

Rohit = 1000 Rs

- ① → New 9000
- ② → Andhra 500

→ **Joins / Complex Queries**

∴ tables are related → slice & dice & join data

→ **Strongly defined / rigid Schema**

each column has a datatype  
each row has a fixed set of columns

**Shortcomings of SQL databases**

→ rigid schema → if data is highly flexible  
E-commerce → 100,000+ product types

Products

id	name	RAM	HD	brand	size	neck collar
						color

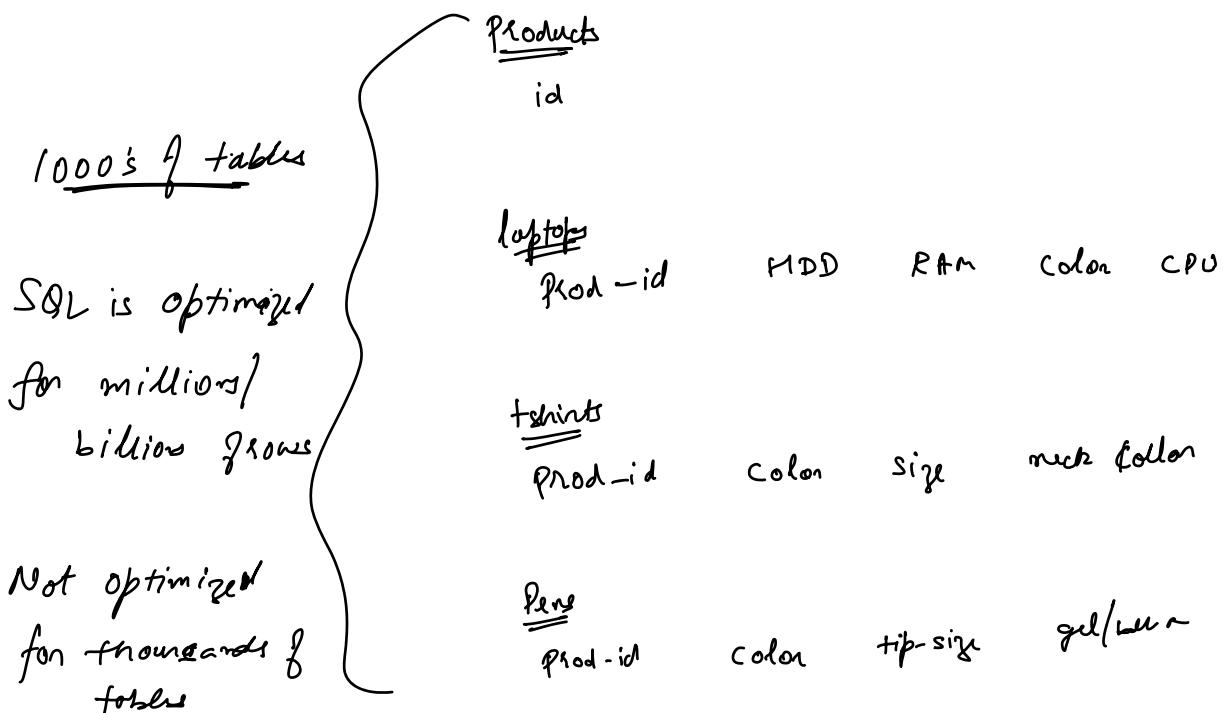
each product has a different set of attribut.

① add all the possible types to the column list

↳ each row will have thousands of  
**null values**

- ↳ Spurious
- ↳ Ambiguity
  - ↳ no data?
  - ↳ is data but we don't know it?
- ↳ Comparisons

② for each product type, I can have a table



If there is a lot of data → single machine cannot handle the load

- ↳ amount of data
- ↳ # reads → multiple replicas
- ↳ # writes

Sharding → only solution

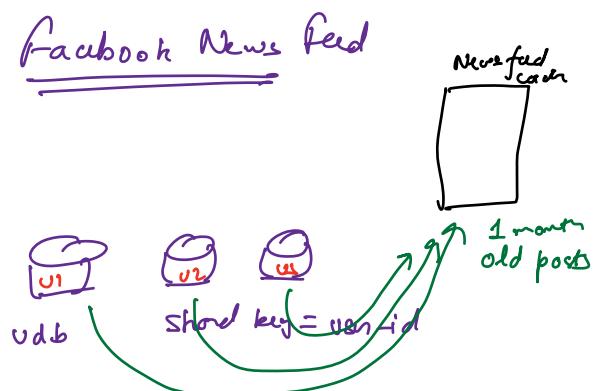
**Sharding nullifies almost all SQL advantages**

{

- ↳ within shard → all good
- ↳ cross-shard
  - ↳ joins → difficult, expensive
  - ↳ relations → not supported
  - ↳ ACID → no longer true across shards

Most popular Relational DBs  
PostgreSQL / SQLite

No official support  
for sharding.



## NoSQL databases

- ↳ key-value
- ↳ document
- ↳ column database
- ↳ graph → much smaller market share
  - ↳ emerging technology

Surreal DB

## Choosing a good sharding key (NoSQL)

chat/messenger

↳ includes SQL

User-id → all message of a user gets a shard

Omkar → Anurag : Hello (m1)

Q: In which shard should we store (m1)

Context: Freq Queries → given a user/conversation  
fetch recent messages

① sender-id      m1 → shard for Omkar

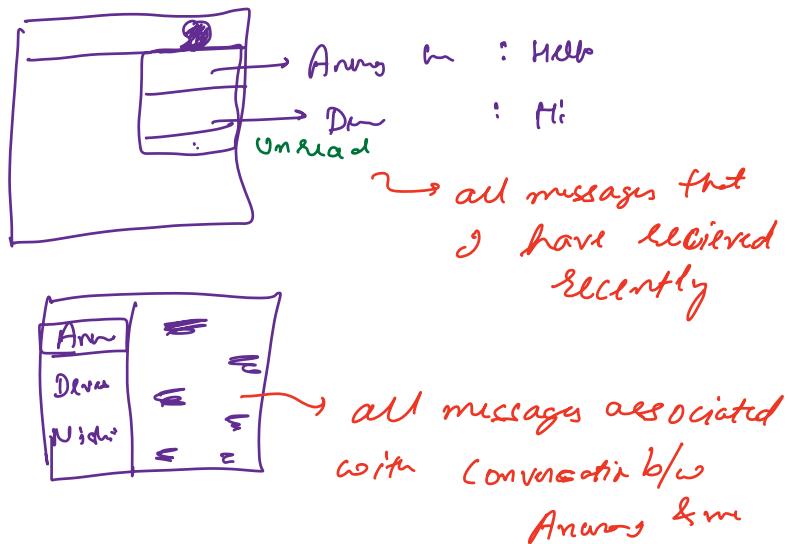
fetching recent messages for  
Anurag → difficult

all sent by me ✓  
all sent to me ✗

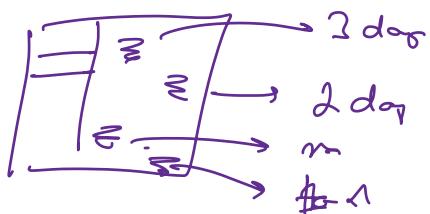
② receiver - id → symmetrically  
Same issue

all sent by me X  
all sent to me ✓

③ conversation - id  
(sender - id, receiver - id)  
fetch all messages b  
this conversation ✓  
find all I sent X  
find all I received X



④ timestamp



⑤ geo - location

Solution shard data still by user-id

but now any message that is sent is stored in both  
↳ sender's shard  
↳ receiver's shard



## ① Banking System

Context ① users are from multiple cities

② Most frequent queries

→ Balance Query (account-id)

→ Fetch transaction history (account-id)

→ List out my accounts (user-id)

→ Make transaction (sending-account-id, receiving-account-id, amount)

① location → Make transaction  
X bad (not really hit two shards) ✓  
query is diff

- users can change cities
  - ↳ copy transaction
- some cities will have significantly ( $1000n$ ) more activity than other cities

② user-id → all queries except Make Transaction  
good only need 1 shard.

→ Make Transaction  
 ↳ save data in both sender & receiver shard.

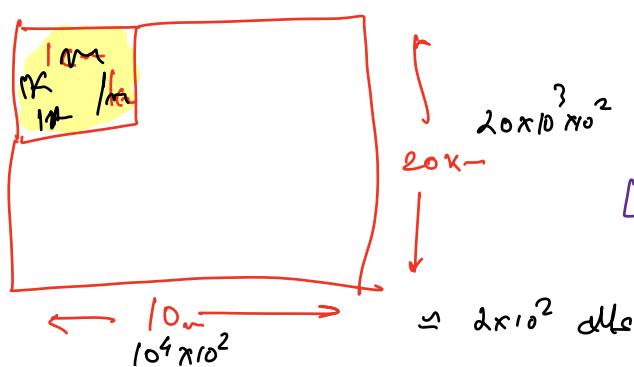
Bangalore → 50,000 riders

③ Ride sharing app

Context → find all nearby cars which are available  
 $(\text{driver-id}, \text{location}, \text{available})$

driver-id → nearby drivers will all belong to different shards

City-id → find nearby cars  
 ↳ drivers can change location



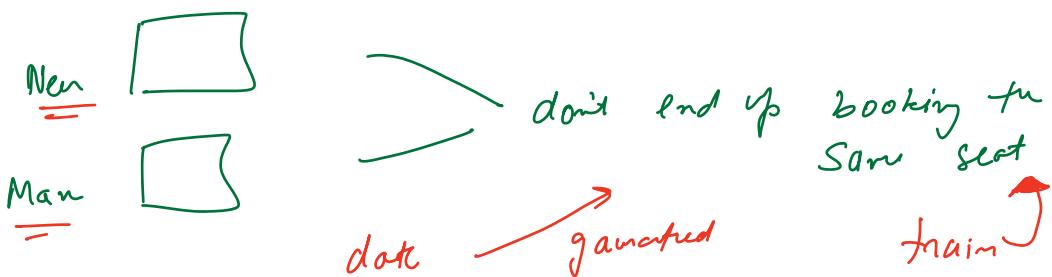
Delhi → Gurgaon  
 ↑ move ten

↳ search for intra-city rides X  
 X drivers must be in same shard

## ⑧ IRCTC

Context → Tatted ticket booking  
 $(\underline{\text{user-id}}, \text{train-id}, \underline{\text{date}}, \underline{\text{source}}, \underline{\text{dest}}, \underline{\text{class}})$

→ 27,000 bookings per minute  
 $\approx 400$  per second.



date → shard for today / tomorrow → overloaded

train-id → Solves our issue completely

1 million \$  
 addition constant → shows only ~~not~~ needed

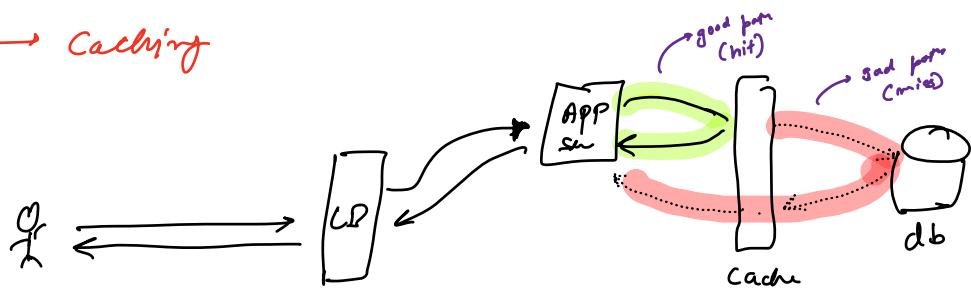
### Summary of choosing sharding key

- load should be evenly distributed across shards
- Most frequent operation/queries/updates should be efficient
- Most freq writes → update least amount of shards  
 $\rightarrow$  better data consistency

→ Minimize (can't always perfectly avoid)  
the data duplicacy.

## Doubts

Caching → Caching



+ve Caching  
-ve Caching

→ response

pragy Solns.com.

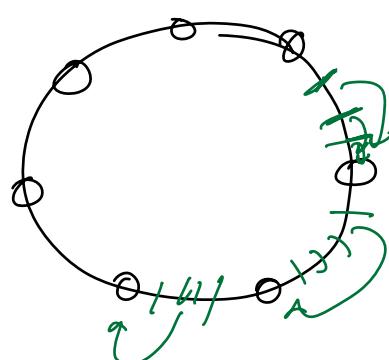
→ is email  
rohan@gmail.com  
registered?

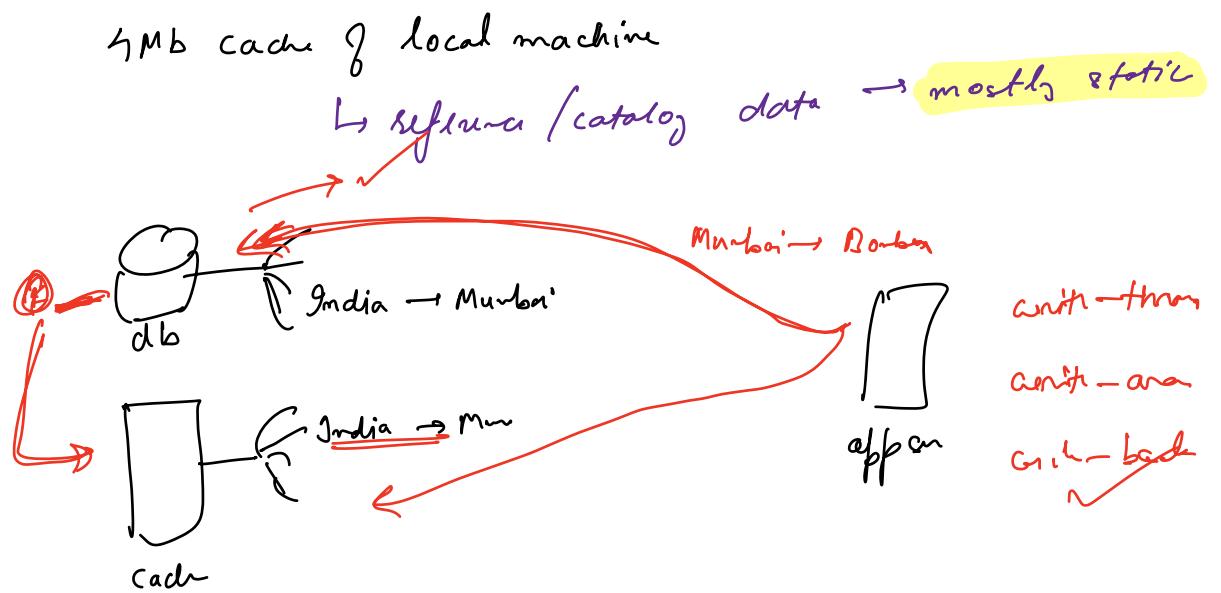
Rohan was right → +ve result from db  
✓ definitely cache it

Rohan was not registered → -ve result

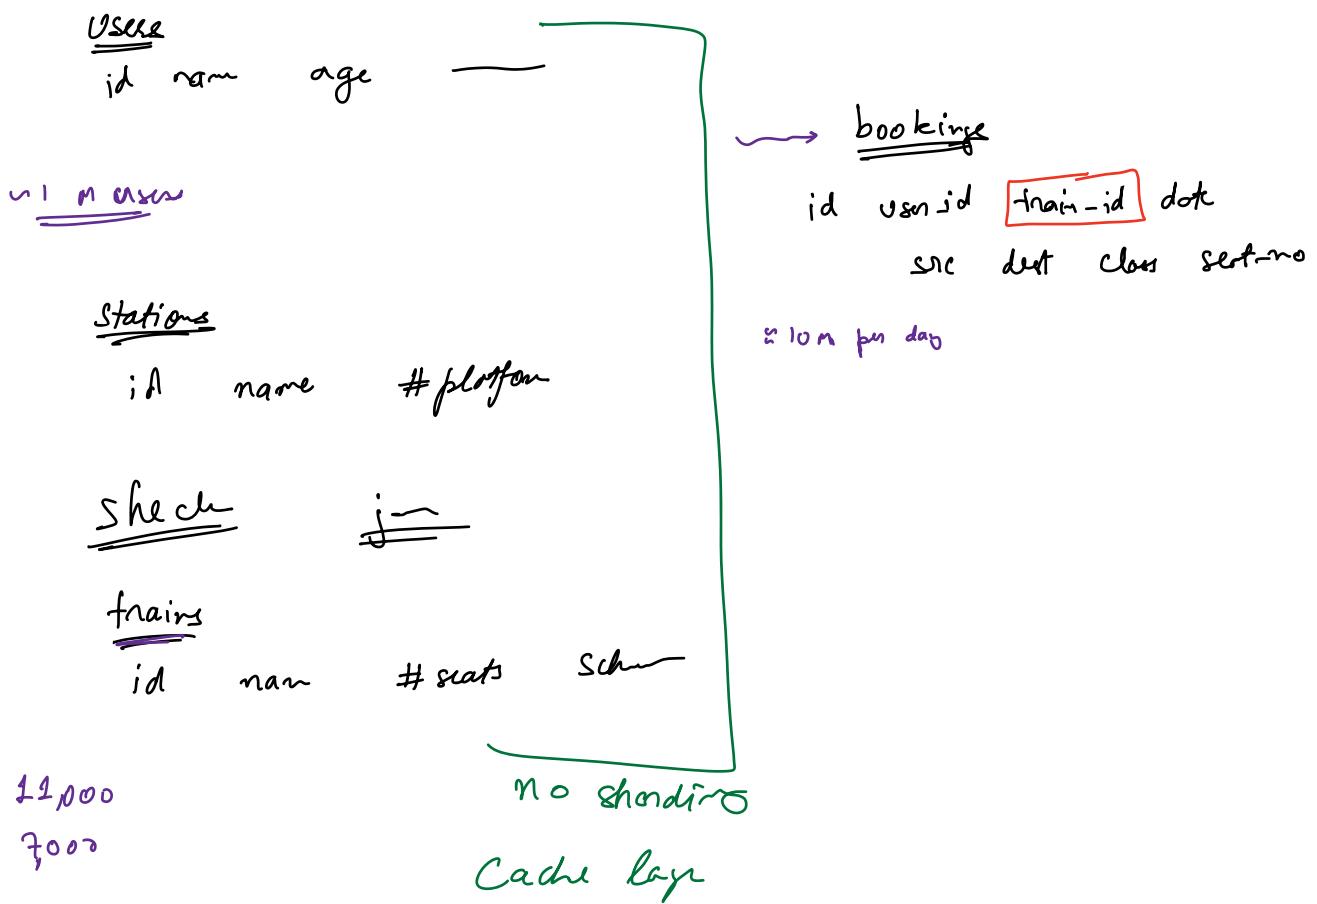
## Consistent Hashing

train-id (11,000 trains)





## IRCTC



group chats → Slack channel 100,000 users

msg → send + receive  
100,000 copies of msg? X

Users  
id name avatar  
user-groups  
id group-

→ sharded by user-id

user-user-chat (DMs)

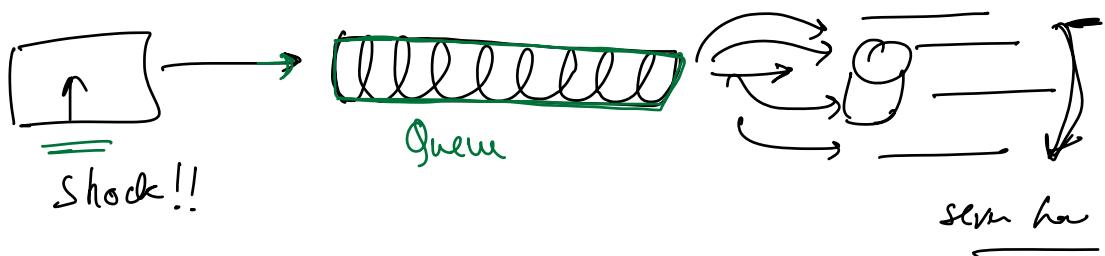
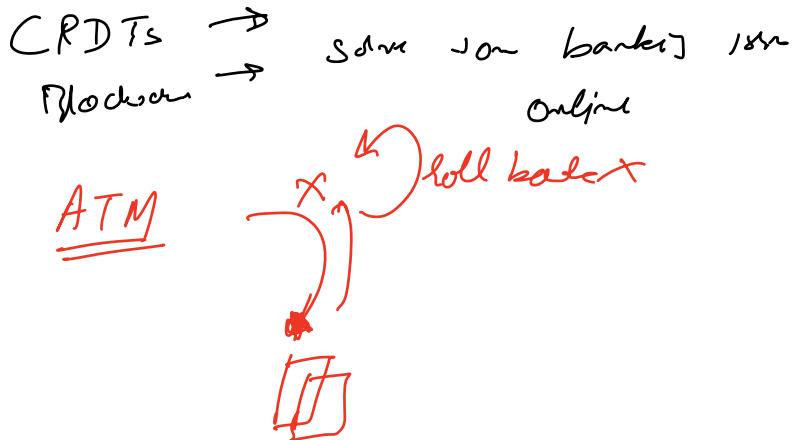
sender-id receiver-id msg timestamp  
sharded by user id  
msg copied  
sender receiver

group-chat  
id sender-id group-id msg timestamp

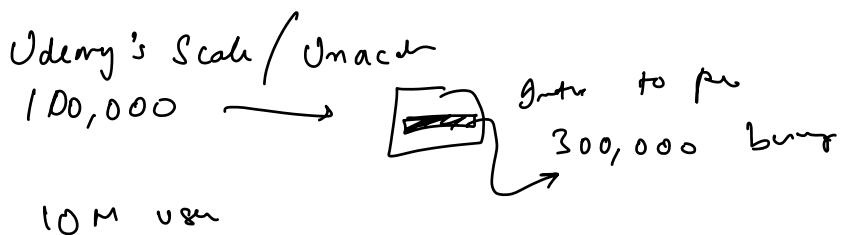
shard on group-id

group-chat-read  
group-id msg-id user-id  
10 200 1

websockets

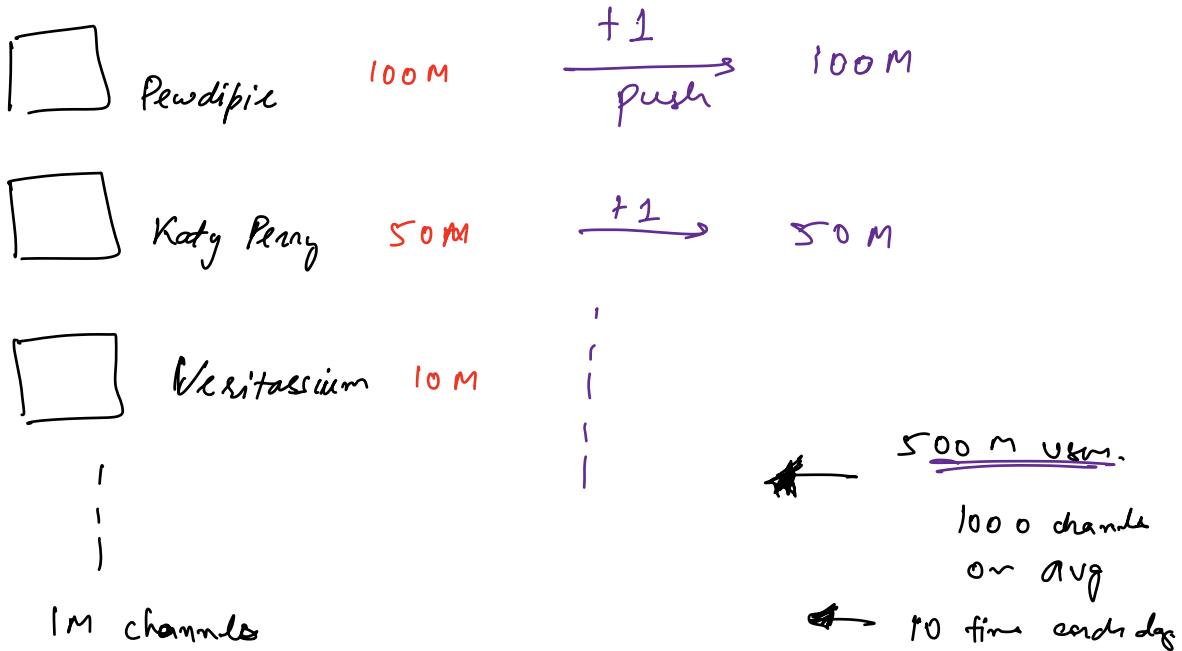


# users = 1,000 → 1500 → split → watch



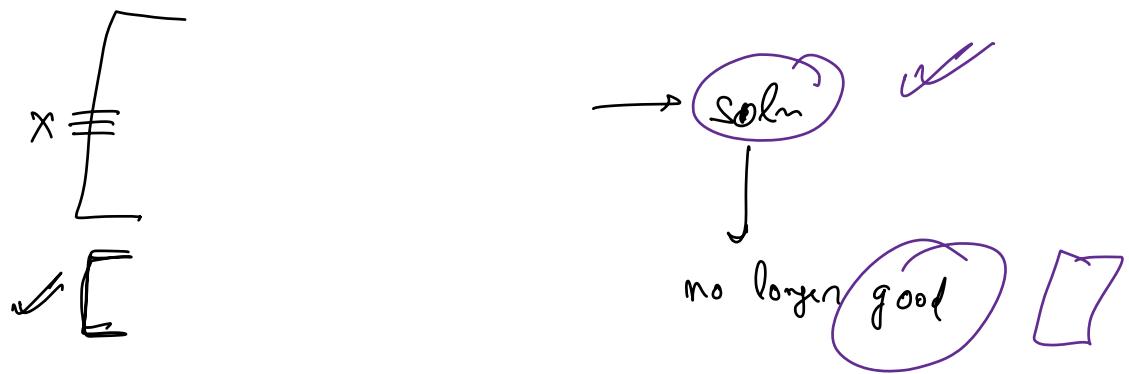
Tanstack  
 → React Query v3  
 Client-Side Cache  
 { Persist Cache in  
 Localstorage / Indexed DB

- SWR
- Non-nalijg back in  
Grager & L

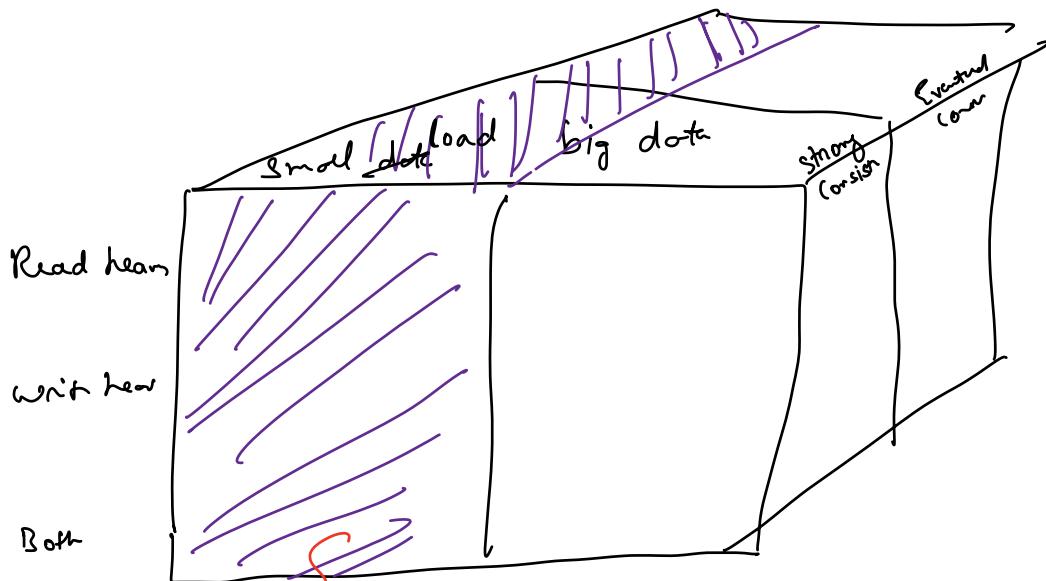


$$\# \text{ events in day} = \left( \begin{array}{c} \# \text{ videos per day} \\ * \# \text{ user subs to perf char} \end{array} \right)$$

$$\# \text{ events per day} = \frac{500 \text{ M}}{\cancel{1000}^{\cancel{10}}} \Rightarrow$$

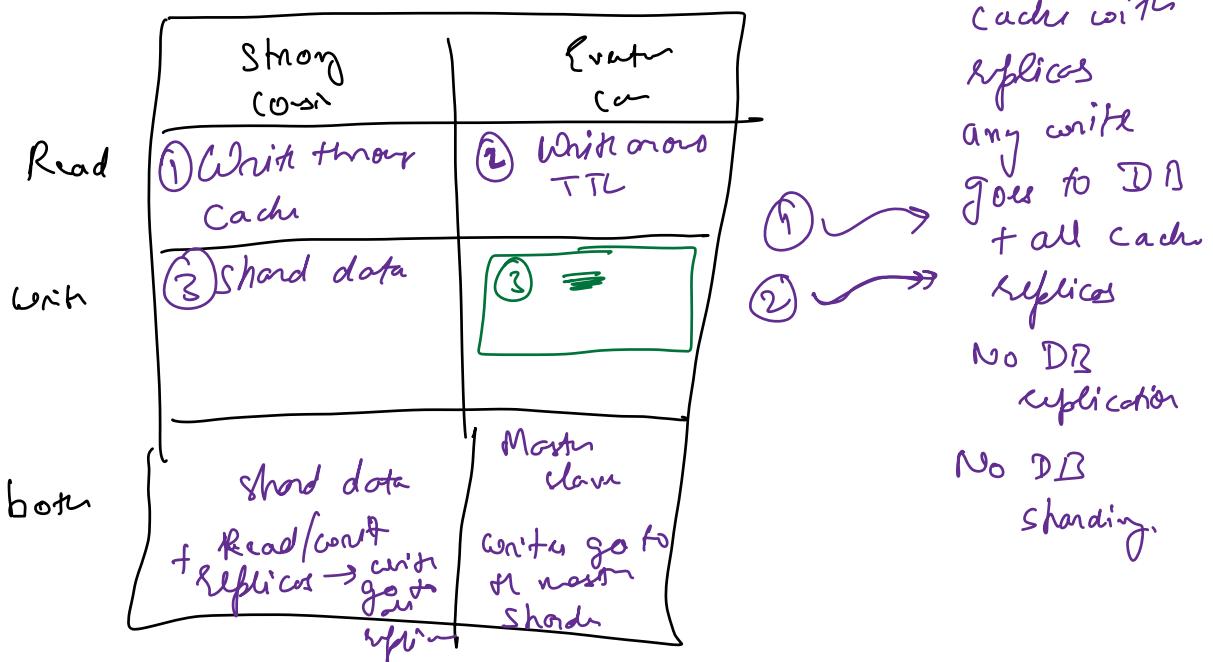


- ① anticipate & pre-emptively plan for it
- ② deal with it → salvage as much as possible



fairial → no sharding  
no replication

## Big Load / Data



⑦ Ensures that any write guarantees full region strong consistency fail in a single shard

Replica → only helps optimize reads & availability

Sharding → helps optimize both reads & writes but not availability.

→ bad if leads to inconsistency  
if not managed

Duplicacy → Redundancy → good ∵ It prevent loss & gives fast reads

Redundancy ← duplicacy