


Agenda


- Quad Trees
- Nearest Neighbor (Uber)

Nearest Neighbor


Circle
↳ very slow



Rectangle
↳ slow




Grid → fast
↳ impossible to
find perfect grid size





Quad Trees

any cell/leaf node will have at most 5 items.



List<Point2D> TL items = items.filter(
 $i \Rightarrow (i.x \leq midX \ \&\& \ i.y \leq midY)$
)



6 items with exactly
 same coordinates!!

Node createQuadTree (point2D TL,
 point2D BR,
 List<point2D> items)

if (items.size() <= 5)

return Node

TL, BR, items

check if overflowing items

$midX = (TL.x + BR.x) / 2$

$midY = (TL.y + BR.y) / 2$

TL items = ---

TR items = ---

BL it = ---

BR item = ---

TLchild = createQT (TL.x, TL.y,

{midX, midY},
 TL items)

TRchild = -- (midX, TL.y,

BR.x, midY,
 TR items)

BL child = ---

BR child = ---

Node parent = new Node()

parent.tl = TLchild




parent.tr = TR---

parent.bl =

parent.br =

} not part

find nearest Locations $(x, y, k \rightarrow 10)$



Node {

* TL $\rightarrow 8$ bytes
 * TR $\rightarrow 8$
 * RL $\rightarrow 8$
 * BR $\rightarrow 8$
 Count $\rightarrow 8$
 fl coords $\rightarrow 16$
 br coords $\rightarrow 16$

}

72 bytes \rightarrow 100 bytes

LeafNode {

List < point 2D, id > 1 km;

neighbor pointers [8] $\rightarrow 8 \times 8$

fl coords $\rightarrow 16$

br coords $\rightarrow 16$

}

100 bytes

In total \rightarrow 100 million restaurants.

\rightarrow location	16 bytes
\rightarrow id	8 bytes
	<hr/> 24 bytes

Total space $100 \times 10^6 \times 24 \text{ bytes} + (\# \text{ nodes}) \times 100 \text{ bytes}$

every leaf node has on avg 1 item.


leaf nodes = 100 million.

parent nodes = 100 million / 4

grandpa nodes = 100 million / 4^2

...

max possible # nodes = 100 million $\left(1 + \frac{1}{4} + \frac{1}{4^2} + \frac{1}{4^3} + \frac{1}{4^4} \dots \right)$



$$\sum_{i=0}^{\infty} r^i = \frac{1}{1-r} \quad |r| < 1$$

max possible # nodes for

$$100 \text{ million places} = 1.33 * 100 \text{ million}$$

$$\frac{1}{1-1/4} = \frac{1}{3/4} = 4/3 \approx 1.33$$

Total space

$$\underline{100} \times \underline{10^6} \times 24 \text{ bytes} + (\underbrace{\# \text{ nodes}}_{1.33 \times 10^2 \times 10^6}) * 100 \text{ bytes}$$

$$10^8 \left(\overset{200}{24 + 133} \right)$$

$$\approx 2 * 10^{10} \text{ bytes}$$

$$= \underline{\underline{20 \text{ GB}}} \rightarrow \underline{\underline{\text{RAM}}}$$

API for Uber \rightarrow Intra-city rides

\downarrow
within

find Nearest Cabs (x, y)

track ride (driver-id)

update location (driver-id, x, y)
new location

Stakeholders

Passenger.

Driver.

\rightarrow asyn job \rightarrow sends
all nearby drivers
a push notification

ACCEPT/REJECT

\downarrow
first driver that accepts
gets the booking.


\rightarrow mark driver as
unavailable
until ride ends

Scale of Uber

large city \rightarrow 50,000 cabs

10 billion bookings per year \approx 30M booking per day.

$\approx \frac{30 \times 10^6}{86400} \approx 400$ bookings per second.



Sharding Key → city-id

↓
upto 50,000 drives
active




Quad Tree for each city.

10:22 → 10:30

User data → Name
Photo
Addr
Phone
Car

Drive →

~ 4M drives
in entire
on Uber.



Optimize Location Updates

50,000 driv

2.5×10^4 req/sec

5M driv

2.5×10^6 req/sec

Update ev
2 sec

decide how freq should driv
update \rightarrow once every 20 sec

if driver's location has changed ^{significantly} in
the last 1 min \rightarrow only then
send location updates.

60 km/hr

$\frac{60 \times 5}{180}$ m/s

16 m/s.

220

2500 req/sec

~ 1000 req/sec


2.5×10^5 req/sec

find Nearest driver

$\xrightarrow{\text{read}}$

Quad Tree


20 times/sec




① Update Quad Tree at specific intervals.

even 1 hour lets up to the tree
↓
1 time / sec

→ grid cannot be dynamic!!!
② as long as the driver remains within the cell there is no need to update driver location.




↓
driver-app →
cell-id, tl, br




if location within tl, br → do nothing
else


send_location_update (cell id, x, y)
↓
old cell id new location

Server responds with new cell-id, tl, br.






Mumbai



drives inside leaf node
changes \rightarrow


Structure of the tree
remains \leftarrow

\rightarrow no nodes are
split or merged!



if we make grid structure static \rightarrow it will
not be able to react
to density changes.


\rightarrow every 3 hours \rightarrow reconfigure the grid structure



drim ≤ 5



> 5

< 5



if no of drivs > 20 \rightarrow split cell

if no of drivs < 10 \rightarrow merge cells





latest loca


Key	Value
driver-id	x, y



Track

Key	Value
drn-id	$[x_1 y_1, x_2 y_2, x_3 y_3]$



Uniqus Quad Tra for each city.






10 sec of down time

draw.io

excalidraw.com



$R+W \geq X$



high consistency

read/write heavy

↳ split into

diff services

if reads don't dep on the heavy writes




↳ heavy sharding

IRCTC

Cascading failure

↳ 1 syst. fails

↓
other syst. to
also



threshold ≤ 5
↳ leaf