

## Agenda

→ Unique Id Generator

→ Rate Limit

→ Garbage Collection

Unique Id generator → millions of ids (no collision)

UUID v4 → 122 <sub>01</sub> 6 bit random number.

1 billion UUIDv4() every second

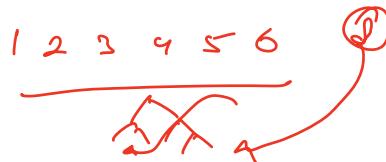
100 years → 50% chance that there will be at least 1 collision.

## Pros

- generated on any system (fully distributed)
- efficient to generate

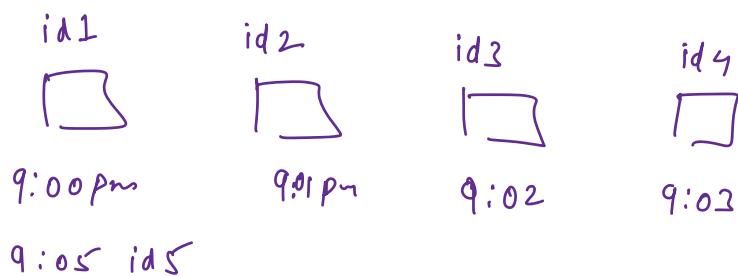
## Cons

- Super long
- Random → not sorted  
Indexing is slower

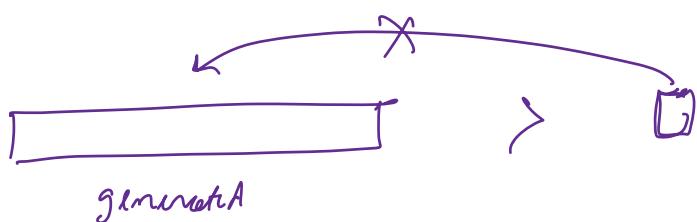


Unique & Incremental id

↳ monotonically increasing.



$\text{id1} < \text{id2} < \text{id3} < \text{id4} < \text{id5}$



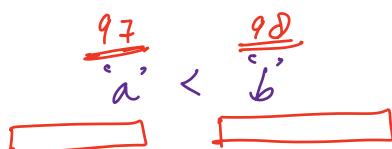
Why monotonically Increasing?

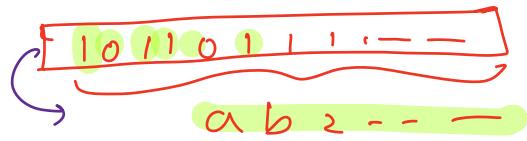
→ Indexing is faster

→ useful for figuring out which event

happened before.

*distributed transactions*





## Count Based

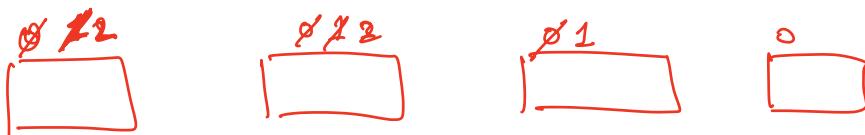
```

int id = 0
int unique-increased-id() {
    return id++;
}

```

— work if there is  
1 system

— many systems



if every system has its own counter

- not increasing.
- not unique

a central system/service that is in charge of ALL id generations.

control API

central Database

### Pros

- meets requirements
- very simple to implement

### Cons

- single point of failure
- latency is high
  - ↳ n/w call for each id
  - ↳ single server will be bottleneck.

a central service  
which is replicated

both read &  
write heavy.

- predictable  $\rightarrow$  security risk.

addhn.india.gov/view/ 0  
1  
2  
|  
3728  
|

addhn.india.gov/view/ 13bc5a - - -  
- - -

$2^{122}$  possible uuids  
 $\approx 10^{36}$

$$\frac{10^{36}}{1.5 \times 10^9} = 10^{27} \text{ ids}$$

Timestamp  $\rightarrow$  every CPU has an internal Quartz clock

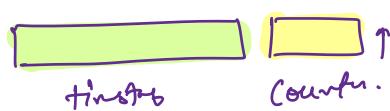
$\hookrightarrow$  System clock  $\rightarrow$  time in milliseconds  $\rightarrow$  integer

(epoch time  $\rightarrow$  # of millisec  
that have elapsed since

1 Jan 1970 - -)

① Single Machine  $\rightarrow$  id() called twice within same  
millisecond

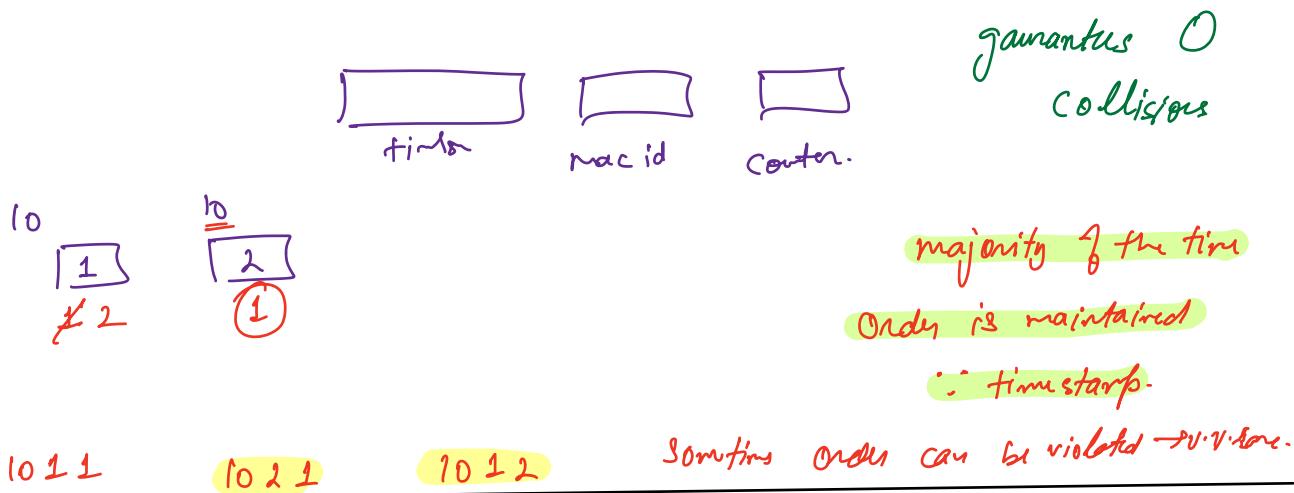
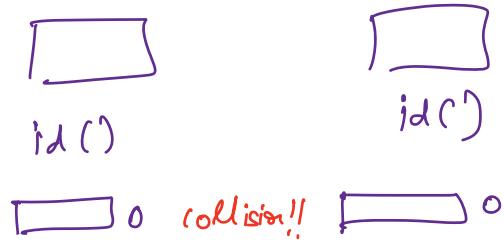
have an additional <sup>thread-safe</sup> counter.



$$\begin{aligned} id1 &= \underline{id()} \leftarrow \text{first} \\ id2 &= \underline{id()} \leftarrow \text{second} \end{aligned}$$

millsec      rest rot = 0  
change

## ② Multiple machines



Achieving truly monotonically increasing unique ids  
in a distributed setting without a central authority  
→ IMPOSSIBLE



Clock synchronization

NTP → it can rewind time!!

$$\text{start\_to} = \text{smst}()$$

~~$\text{end\_to}$~~

$$\text{end\_to} = \text{end\_at}()$$

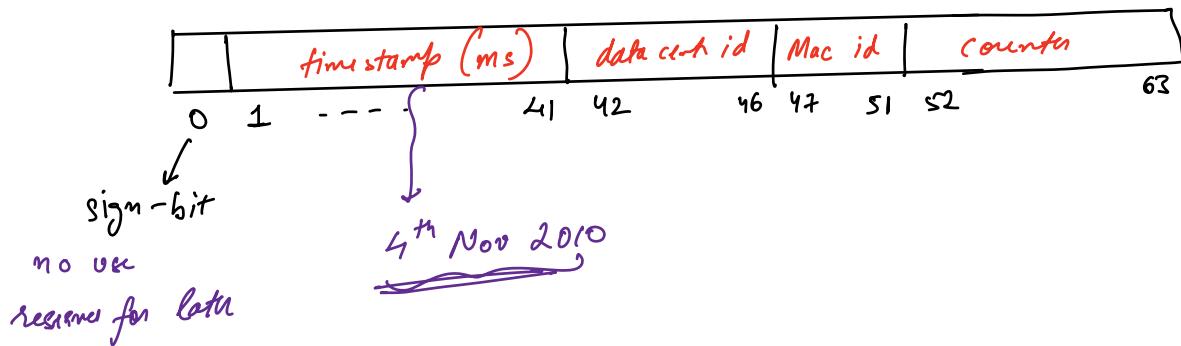
$$\text{elapsed\_to} = \text{end\_at}()$$

'  
— vc

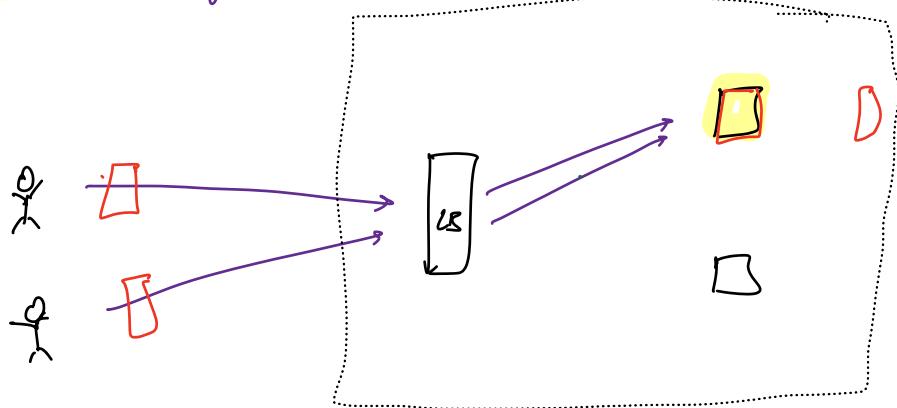
Snowflake IDs — Twitter

Instagram / Discord

64 bit integer



**Rate Limiter** → any user can only make  $k$  requests / seconds



### Reasons

- prevent DDoS attacks
- avoid abuse
- ensure fair usage

Each user has a unique ip address or some other id.

Design rate limits per user.

User hash map       $\text{map} < \text{User-id}, \underline{\hspace{2cm}} \rightarrow$

~~Store count~~

~~map < user-id, {int, first-call}~~

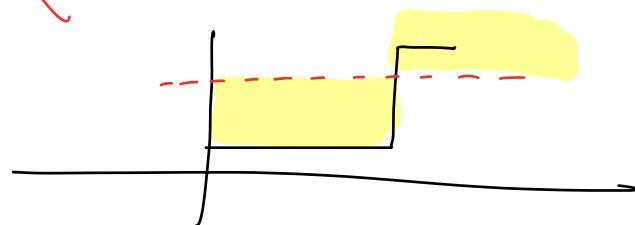
~~void handle-request (req) {~~

~~now = count-tl()~~

~~diff = now - first-call~~

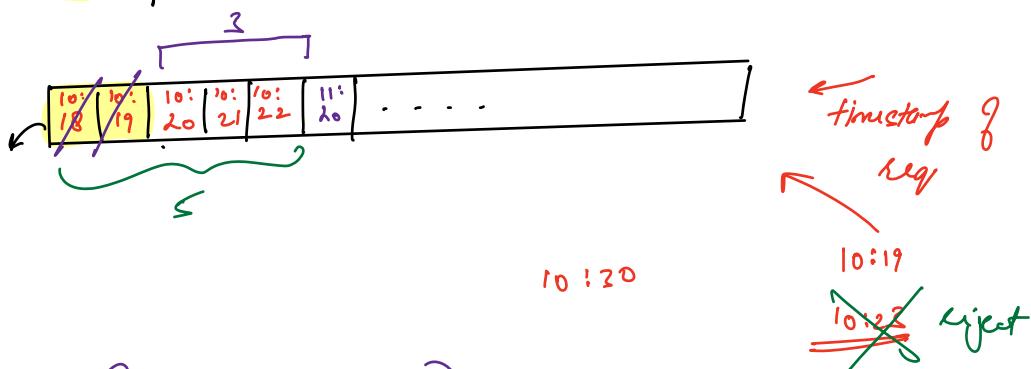
~~count / diff → rotation limit } avg rate  
or violations }~~

*reject()*  
*process (req)*



constraint : 5 req/hour

Request Queue per user



$(10:23 - 1 \text{ hour}) \rightarrow$  any req older than  
this can be zero  
from queue  
current req  
timestep  
 $= 9:23$

$$\cancel{10:30} - 1 \text{ hour} = 9:30$$

$$\cancel{\checkmark 11:20} - 1 \text{ hour} = \cancel{10:20}$$

## Challenges

→ High memory → maintain a queue for each user.

→ High latency for certain requests 50K req/hour



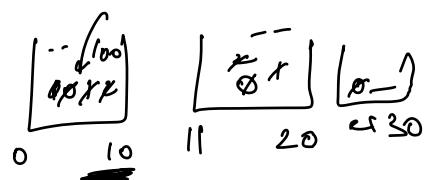
## Bucketing time & use counter

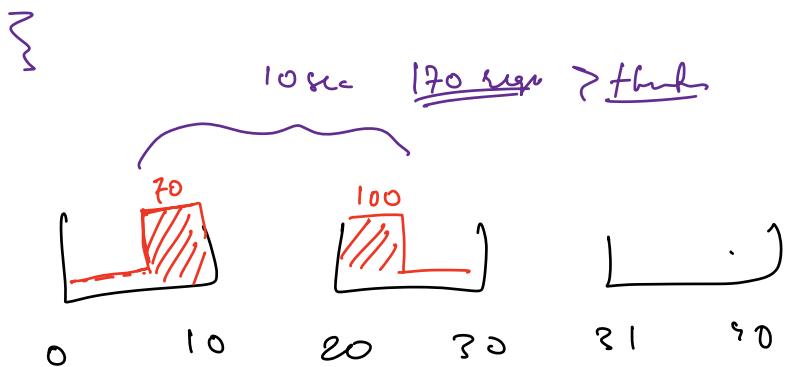
100 req/10 seconds

$$\text{start-time} = \text{now}()$$

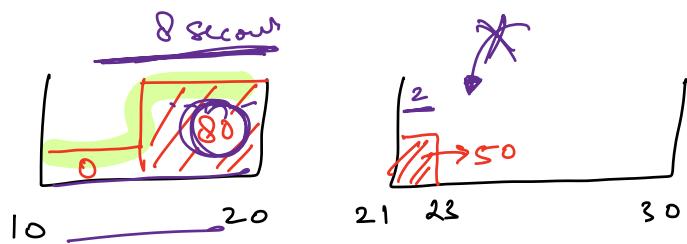
Count = 0

```
void handle-request() {
    current_time = now()
    if (current_time - start_time > 10) {
        count += 1
        if (count > 100)
            reject
        else
            accept
    }
}
```





approximate    (Sliding Window Counter)



8sec/s     $50 + 64 \geq 100$

(64 rps)

map < user\_id, pair < int, int > >

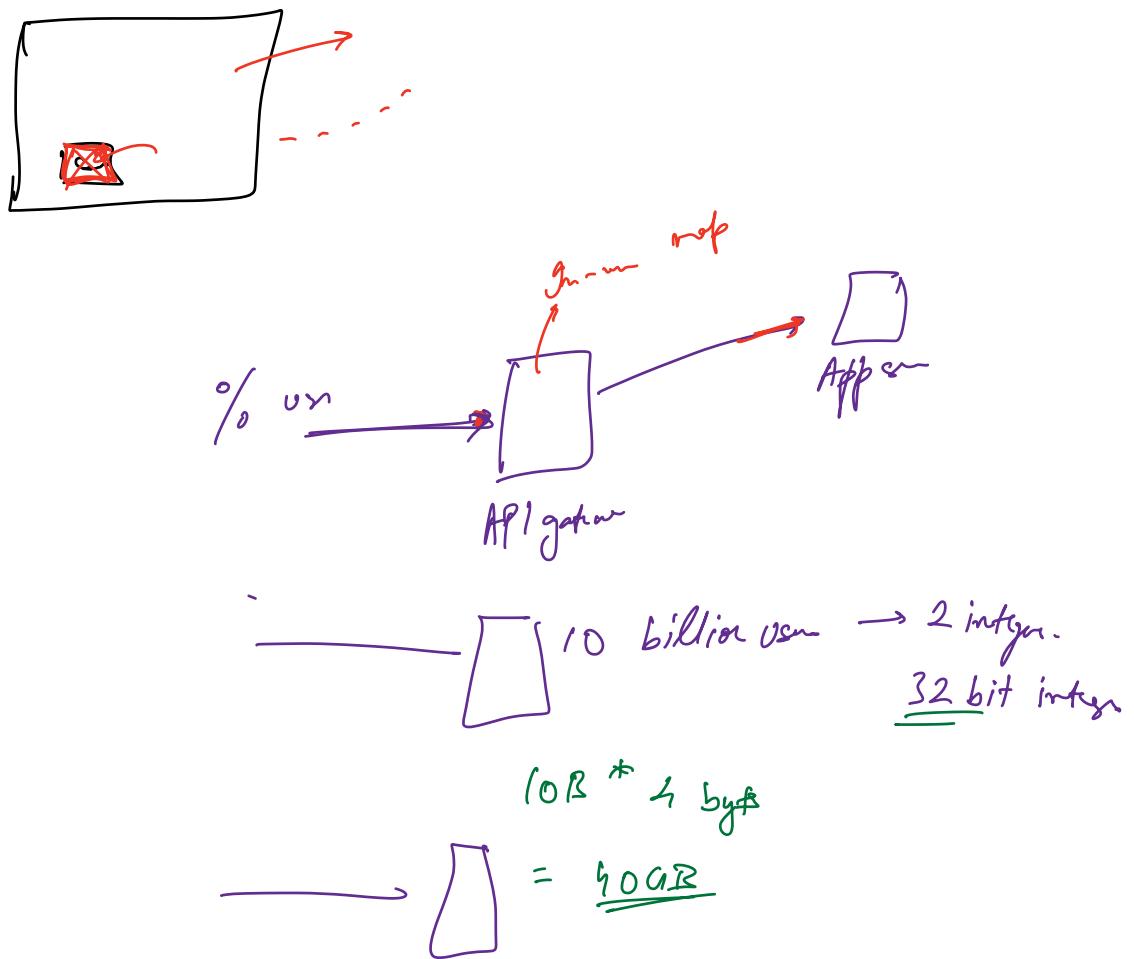
last bucket count  
can back count.

Token

map < user\_id, int >     $\oplus$   $\rightarrow$  token count

7Ps / week

10:40 → 10:50



## Garbage Collection

Why?

if we allocate memory & never de-allocate it  
Memory Leak

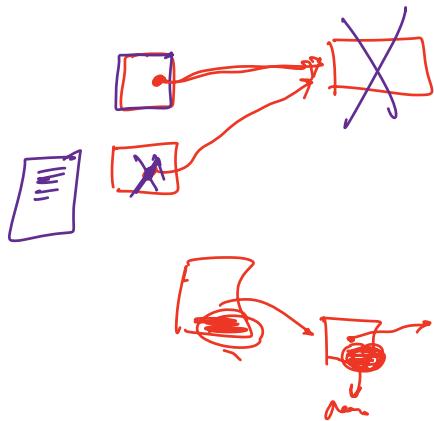
### ① Manual Memory Management (C)

↳ 80% bugs are due to MM

↳ leaks

↳ dangling pointer → points to memory that does not exist

↳ double free



### ② Program to do memory management for us

Garbage Collector.

### ③ Rust → Borrow Checker

## Common Issues with GC

→ overhead

→ Some GC algo can lead to unpredictable pauses!!

## Functional Requirements

- ① system finds & removes objects that are no longer in use
- ② correctness → it should not release any object that is in-use

## Non-functional Requirements

- ① High performance / low overhead
- ② Predictable

---

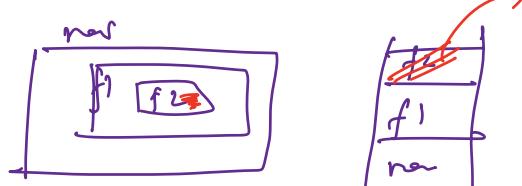
## Stack Memory & Heap Memory

Call stack for each thread

- ↳ functions are currently running
- ↳ variables have been declared to be in scope.

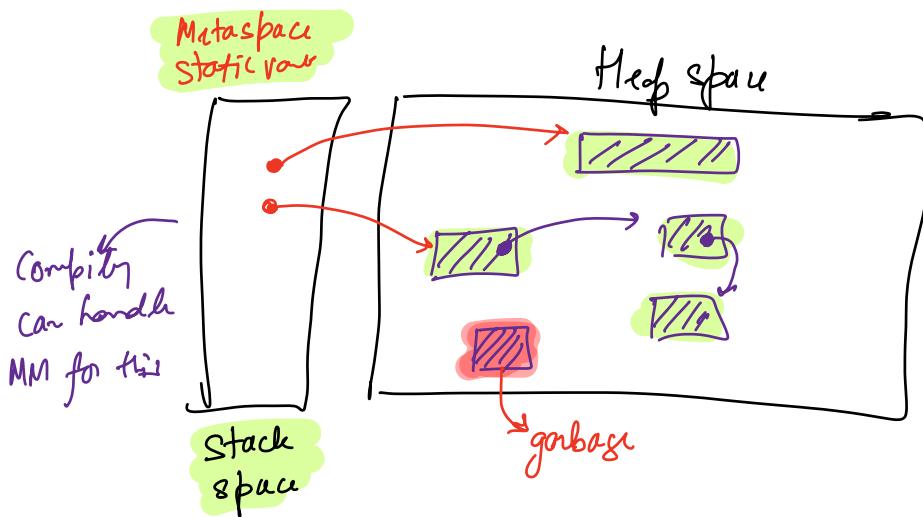
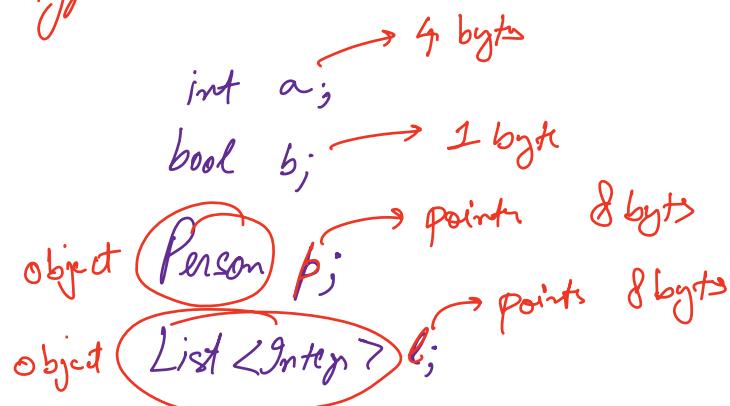
any primitive variables / any static variables

int/pool/-



all variables  $\rightarrow$  stack

↓  
Primitive type



```
:f(a){  
    p(a);  
}  
a=10;
```

## Mark & Sweep

① stops the world (pauses any mutator threads) in your program

② Start from (stack + static + metaspace)

③ DFS → marks all reachable objects

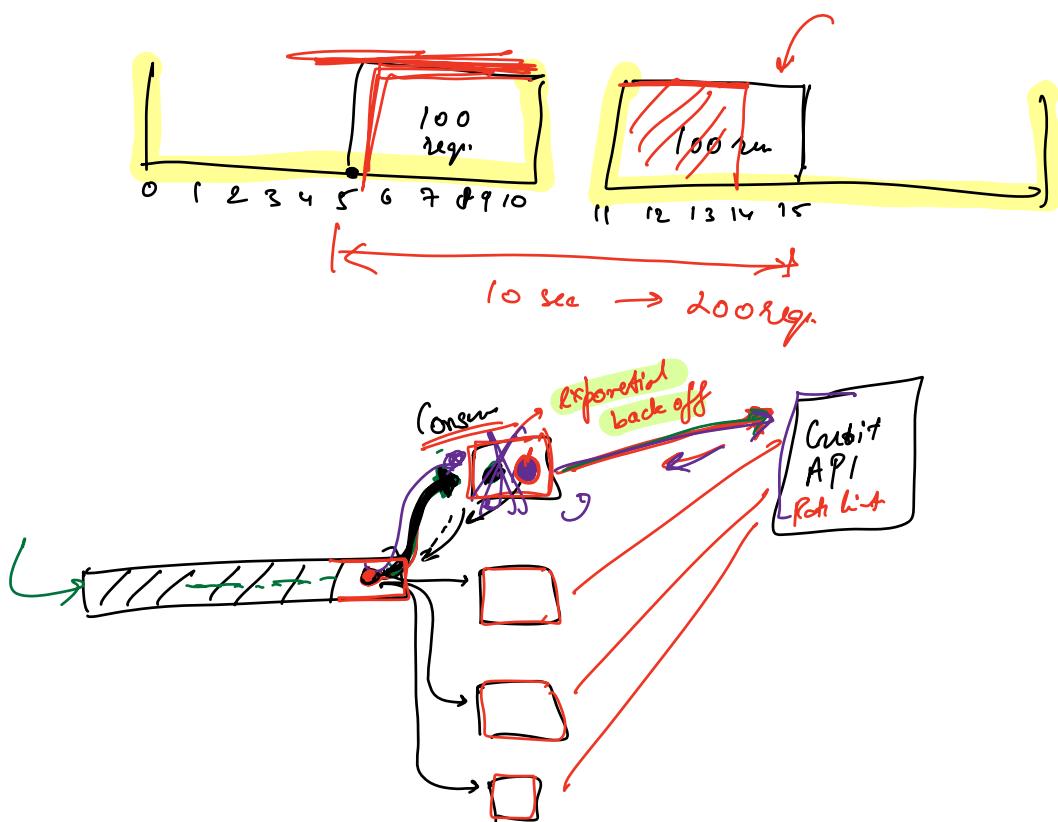
5 seconds

3

Swap entire heap & remove unmarked objects

## ⑤ Defragment memory

Generation GC, Reference Counting, Tri-color marking scheme, O(1)



```

void consumer() {
    check disk
    msg. ← get ←
    msg. ← count →
    write to disk —
    Comm. after
    = { Process
    new for disk
}

```

	Arunraj	Kavithik	Ashwini	Rishabh
Age	25	30	20	26
Height	6'	12'	5"	5'

Print Index on Age

