

Agenda

- ✓ → Problem Statement
- ✓ → Minimal Viable Product
- ✓ → Scale Estimation
 - Design Goals
 - API
 - System Design

Design / sharding / caching /

database

database schema

Problem Statement

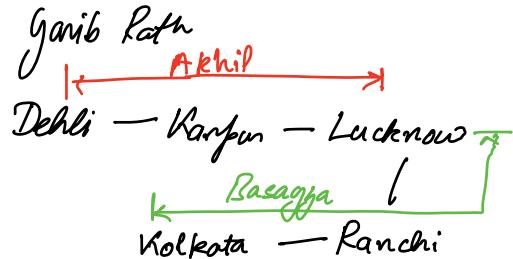
IRCTC → Online ticket booking service for Indian Railways → Transport backbone of India

Tatkal Emergency → opens 1 day before at 10am — huge traffic spike

most populous country.

Minimal Viable Product (v0)

- login & registration
- search for trains
- train details
 - planned-schedule
 - stations & day+time
- seat availability (class)
- booking - passenger details
- payment
- seat allotment
- cancellation (v1)
- station wise seat breakdown (v1)
- seat preferences (v1)
- RAC (v2) → PNR
- notifications (v2)
- Integration with other services (accommodation booking) (v3)



Scale Estimation

end goal → estimate # queries / sec
↳ estimate data size

India Population → 1.5B

users → 1% of Indians are registered on IRCTC

$$\frac{1.5B}{100} = \frac{1.5 \times 10^9}{100} M = \underline{15M} \text{ users}$$

$$5\% = \underline{75M \text{ users}}$$

Actual # users ≈ 40M

MAU = all user base → 40M

DAU = 10% of entire user base = 4M users

Estimate # bookings per day = 1-5M per day not every one might be booking (20-80)
20% of 4M

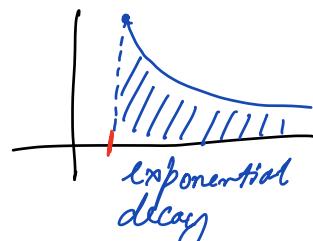
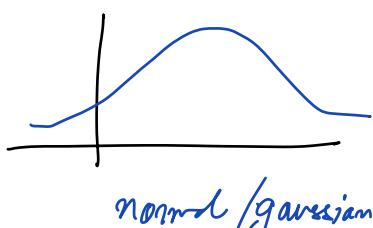
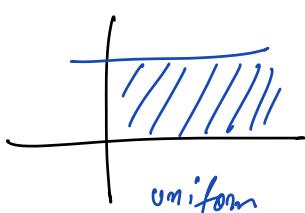
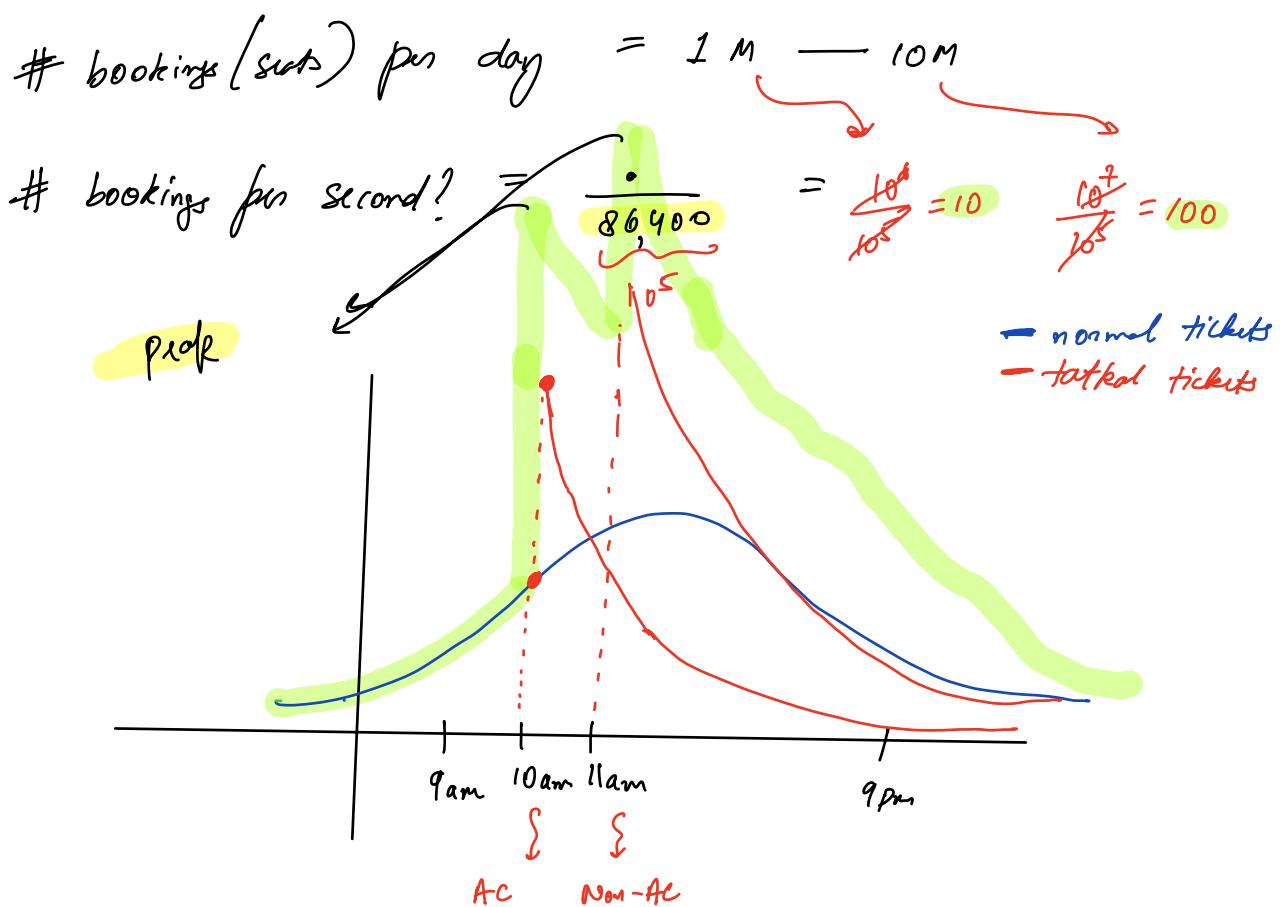
Nikhil → # trains = 10,000 → 10^4

assume 100% occupancy

$$\begin{aligned} \text{approx. # seats per train} &\rightarrow \underbrace{\# \text{ coaches per train}}_{15} * \\ &\quad \underbrace{\# \text{ seats per coach}}_{72} \\ &= 15 * 72 = 1080 \xrightarrow{\sim} 10^3 \\ &\quad 9 \text{ segments/coach} \\ &\equiv 11 \quad \left. \right] 8 \text{ seats per segment} \end{aligned}$$

$$\begin{aligned} \text{approx \# seats booked per day} &= 10^4 * 10^3 \\ &= 10^7 \approx 10 \text{ million.} \end{aligned}$$

tickets < # seats each ticket on avg books } 2 seats
 $\therefore 1 \text{ ticket can demand } \frac{10^7}{2} \text{ seats}$

$$\frac{10^7}{2} \Rightarrow \underline{\underline{5M}}$$


✓

(law of large numbers)

(Zipf's law)

many independent distributions

they may add up to gaussian.

Actual scale

↳ # users $\approx 40M$
 ↳ # stations $\approx 7K$
 ↳ # trains $\approx 13K$ (passenger) + 7K (freight)

↳ # tickets/day $\approx 1.1M$ (avg for 2021-2022)

↳ spikes during holiday season

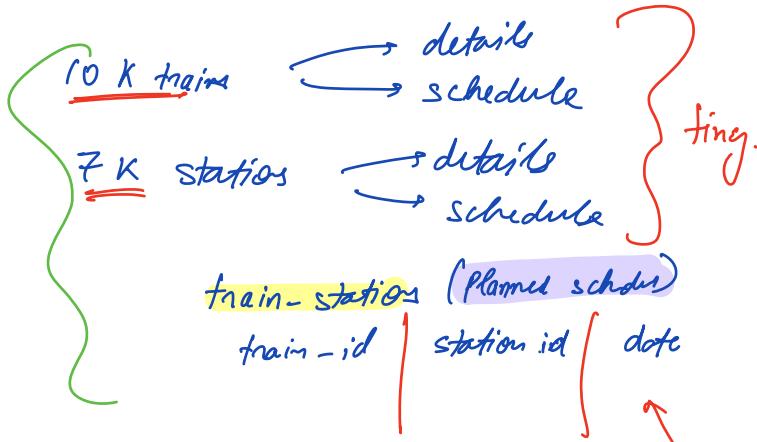
$$\# \text{ tickets/min (peak)} = 26,458 \text{ tickets/min}$$

On Mon 2020

$$QPS \Rightarrow \frac{26K}{60} = 430 \text{ bookings/sec}$$

Amount of data

more or less static



$10k \times 20 \approx \frac{1}{\text{tiny}} \text{ lookups}$

Booking ↑ charges

8	8	3	4	3	1	8	4
id	user-id	train-id	sic	dist	class	created-at	seat-id
PK	FK	FK					

Booking-passenger

8	50	1	1	1	1	1	30
booking-id	passenger-name	age	meal	boat-pref	Doc-type	Doc-id	

size of row = 50 bytes

booking / day = 1M

data per day = 50M bytes = 50MB

$$\begin{aligned}
 20 \text{ years} &= 50M \times 20 \times 365 \approx 365 \text{ TB bytes} \\
 &\approx 365 \text{ GB} \\
 &= 500 \text{ GB} \quad (\text{total data})
 \end{aligned}$$

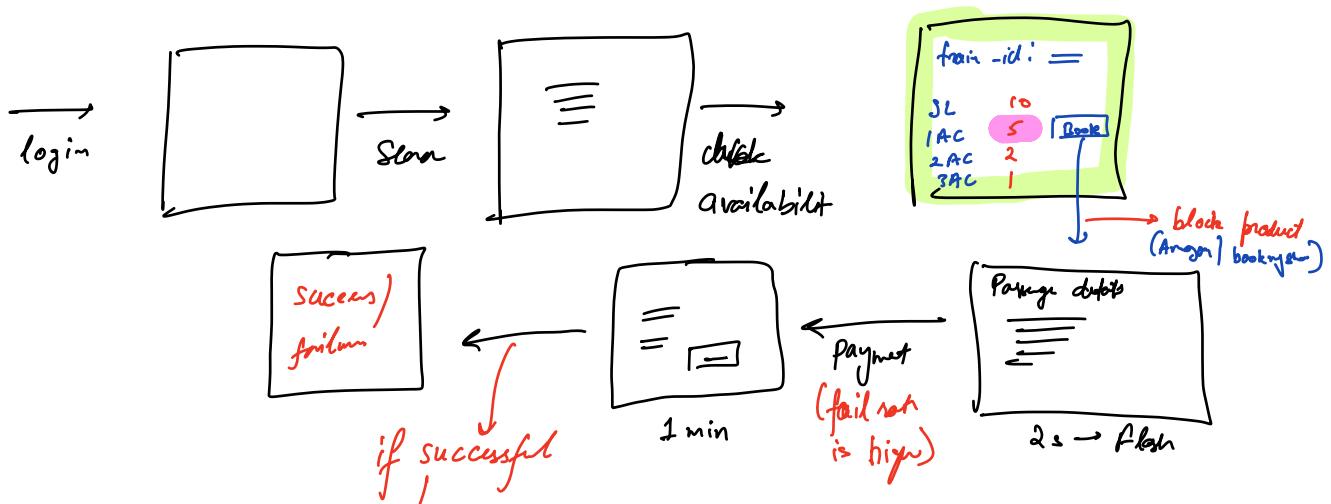
Do we need sharding

∴ of data size? No!!

∴ of load → will see

Design Goals (Non-functional requirements)

- ✓ - login & registration → support microservice
- ✓ - search for trains → availability over consistency.
- ✓ - train details
 - planned-schedule
 - stations & day+time
- ✓ - seat availability (class) → read/write once a month
- ✓ - booking - passenger details → read/write
- ✓ - payment - independent service → consistency/availability
- seat allotment → consistency/availability
read/write
- cancellation (v1) → no read/write (as input)
- station wise seat breakdown (v1)
- seat preferences (v2) → availability/consistency
read/write
both! → check
block



Latency stat

This is a significant time gap b/w seeing the availability & seat allotment.

Design

- ① List of trains
 - ② List of stations
 - ③ Search train based on [Source, dest, date]
 - ④ Train details
 - └ number
 - └ planned schedule
- System - 1
- ③, ④ → read heavy
most by static
pref. availability.
little data
relational

SQL + Redis Cache

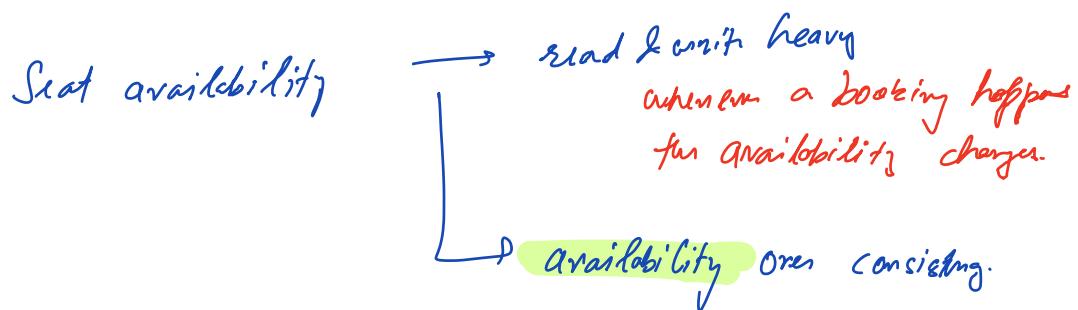
↳ invalidation
(units → streams)

trains				
id	name	number	details	---
1	ganib	100	—	

stations				
id	name	num-platforms	category	
1	Dahli			
2	Luck			

Planned-schedule					
train-id	index	station-id	day-of-week	arrival-time	departure-time
1	0	1			
1	1	2			

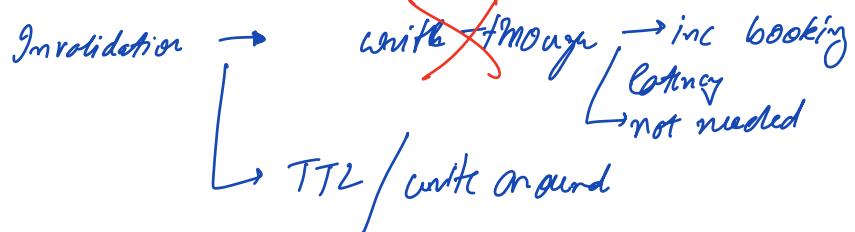
1 2 3



System - 2

DB → SQL

Redis cache



10:48 → 11:00 pm

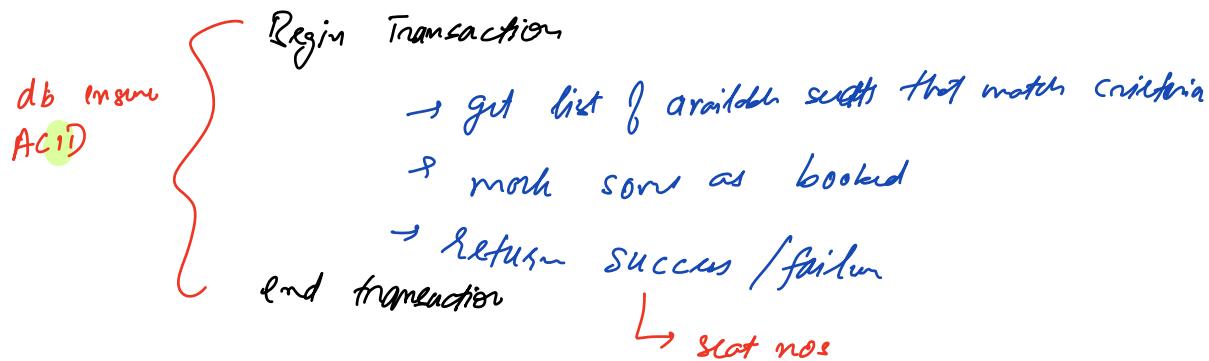
Seat Allotment / Ticket Booking

(train-id, vrn-id, class, source, dest, date,
#seats, berth-type for each seat, meal for each seat)

- ① ensure seat is available → seat-availability-store
X cannot be used : ∵ can
not immediat consistency
- ② mark seat as booked.



both can be done in a single query/transaction.



400 bookings/sec at peak → a single server can't handle this

dealing with load

cache → reads but not writes

replicate → reads but not writes

shard → Only way to handle high write throughput.

Sharding - key

→ date of journey load is evenly distributed.



train-id

every booking depends on train & every train is independent.

bookings for any date of journey that has passed → move to a diff db at end of day.

CRON Job.



↳ Total live data → amount of active booking
(today → 2 months in future)

1M booking /day
60M bookings (active)
50 bytes / booking
= 3GB of data (tiny!!!)

S. Sudarshan → Head of CSE at IIT-S

↳ god of databases

1ms → 500ms for a query to run in db

1000 units/sec → 2 units/sec

Database Schema

↳ batch program
↳ booking same seat over non-overlapping stations.
↳ check which seats are available

bookings
show booked seats
 $(\sum - \text{booked seats}) \rightarrow \text{available seats}$

pre-book seat-train combinations.

Bookings							
8	8	3	3	2	3	4	boarding date
id	user-id	train-id	src	dst	class	created-at	seat-id
PK	FK	FK					
Shatabdi:					1	today	
Shatabdi:					2		
Shatabdi:					:		
Shatabdi:					1000	today	
Shatabdi:					1	next task	
Shatabdi:					2		
Shatabdi:					:		
Rajdhani:					1000	next task	
Rajdhani:					1		
Rajdhani:					2		
Rajdhani:					:		
Rajdhani:					1000		

- ① explode db to pre-emptively create seat-train combinations for next 60 days
- ② user-id is initially blank
- ③ if seat gets booked → user-id gets filled

update bookings

set usn-id = req.usn-id
src = :
dest = :

→ mark slot as booked
from src to dest

where

train-id = req.train-id

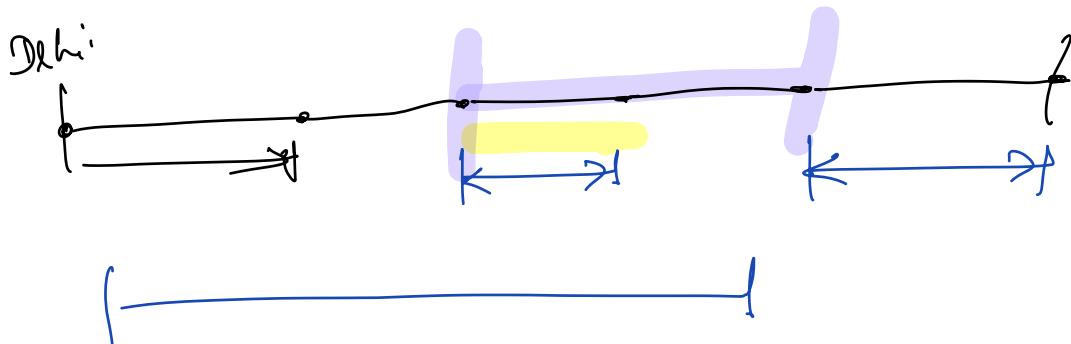
Slot is blocked for
entire journey! (bad)

date = :
class = :

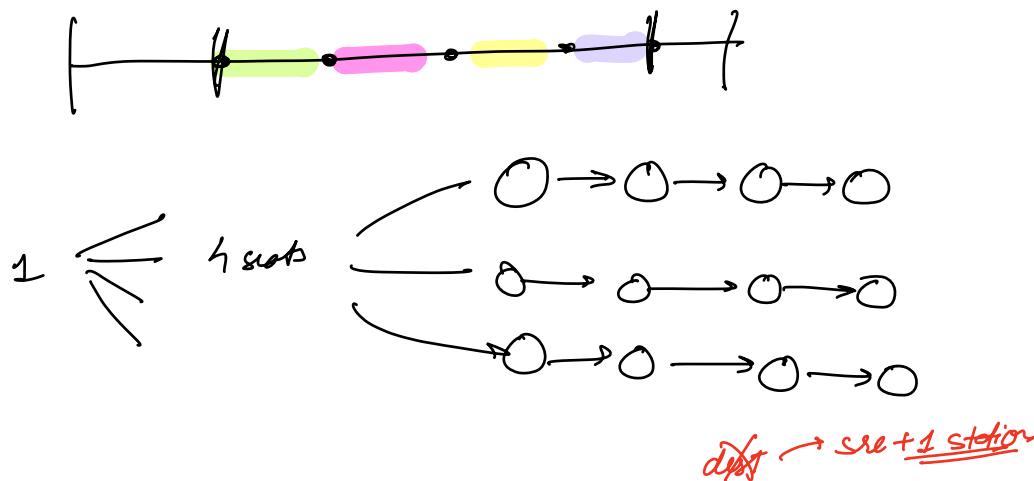
usn-id is NULL → seat is currently
occupied.

Order by select (1 case berth-type = req.bt
0 else) desc

limit req.seat-count



each row in bookings-table only represents journey from 1 station to the immediate next station.



id	user-id	train-id	date-of-journey	src	seat-id	berth-type	class
----	---------	----------	-----------------	-----	---------	------------	-------

10^4 10^2 60

propose for every train, for every seat, for each date,
for each station!!!

~~for #station for train on avg~~
20

$$\text{Size of total} = 10^4 * 10^3 * 60 * 20 = 1.2 \times 10^{10} \text{ entries}$$

each entry is 100 bytes \Rightarrow total size = 1.2×10^{12} bytes

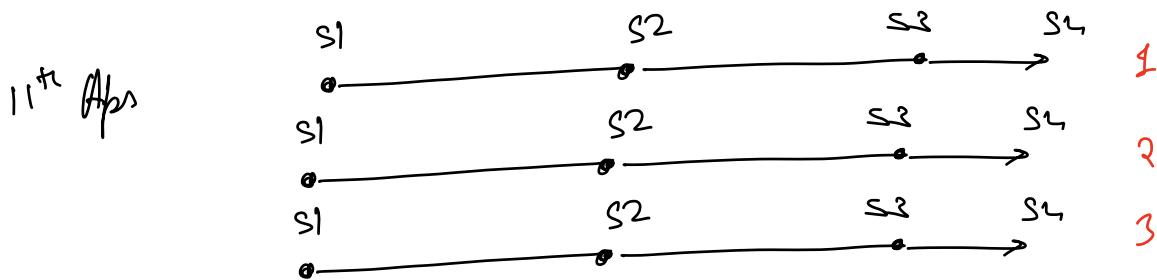
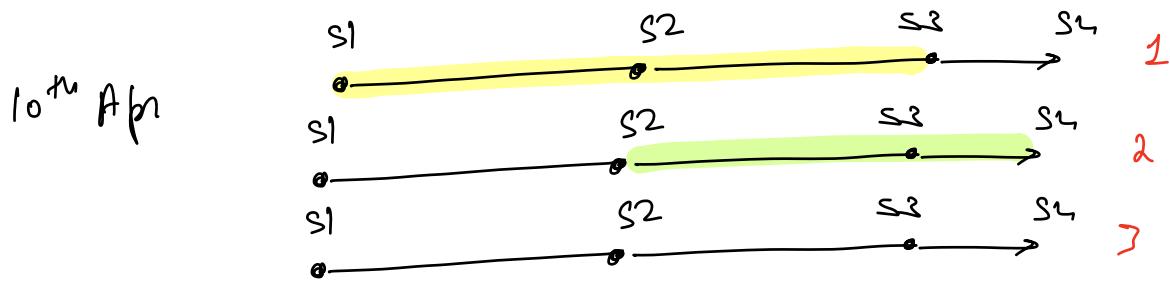
$$= 1.2 \text{ TB}$$



each train has 3 seats $\frac{1}{1} \frac{2}{2} \frac{3}{3}$
 two days → 10, n^4



id	vsn-id	train-id	date of journey	src	seat-id	<u>berth-type</u>	class
1	v	T_1	10	S_1	1		
2		T_1	10	S_1	2		
3		T_1	10	S_1	3		
4	v	T_1	10	S_2	1		
5	n	T_1	10	S_2	2		
6		T_1	10	S_2	3		
7		T_1	10	S_3	1		
8	n	T_1	10	S_3	2		
9		T_1	10	S_3	3		
10		T_1	11	S_1	1		
11		T_1	11	S_1	2		
12		T_1	11	S_1	3		
13		T_1	11	S_2	1		
14		T_1	11	S_2	2		
15		T_1	11	S_2	3		
16		T_1	11	S_3	1		
17		T_1	11	S_3	2		
18		T_1	11	S_3	3		

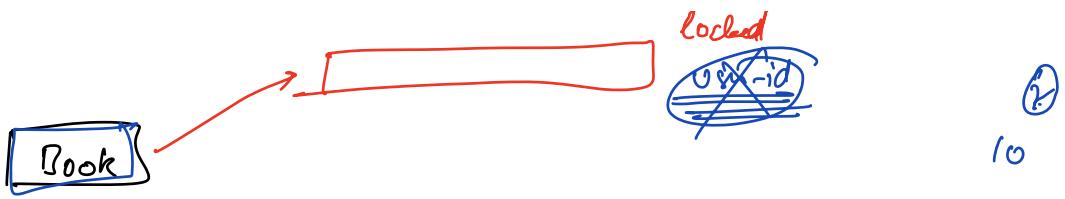


Vorab → 10^{th} 1 seat $\begin{matrix} S1 \xrightarrow{\text{to}} S2 \\ \downarrow \quad \downarrow \\ S1 & , & S2 \end{matrix}$

Nurzij → 10^{th} 1 seat $S2 \rightarrow S3$

$S2 \rightarrow S3$
 $S2 \rightarrow S1$

- Birth prefen
- gender
- VIP
- Quotas (woman, elderly, veterans, VIP, staff)



72 props for begin → 7 tons

bogie → 1000 tons.

- ① List nouns
 - ② identify noun entities
 - ③ relations
-

Read heavy → Caching / replication

Write heavy → Sharding

Read & write heavy

↳ break down into components

↳ R+W heavy → sharding.

Master - Slave

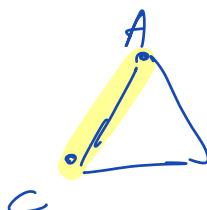
- ① write \rightarrow leader
read \rightarrow any slave } No consistency
- ② write \rightarrow leader + 1 slave
read \rightarrow any slave } eventual consistency,
high availability
- ③ write \rightarrow leader + all slaves
read \rightarrow any slave } immediate cons
low availability
reads fast writes slow
- ④ write \rightarrow majority servers
read \rightarrow majority servers } Quorum
immediate cons
somewhat available
reads & writes slow

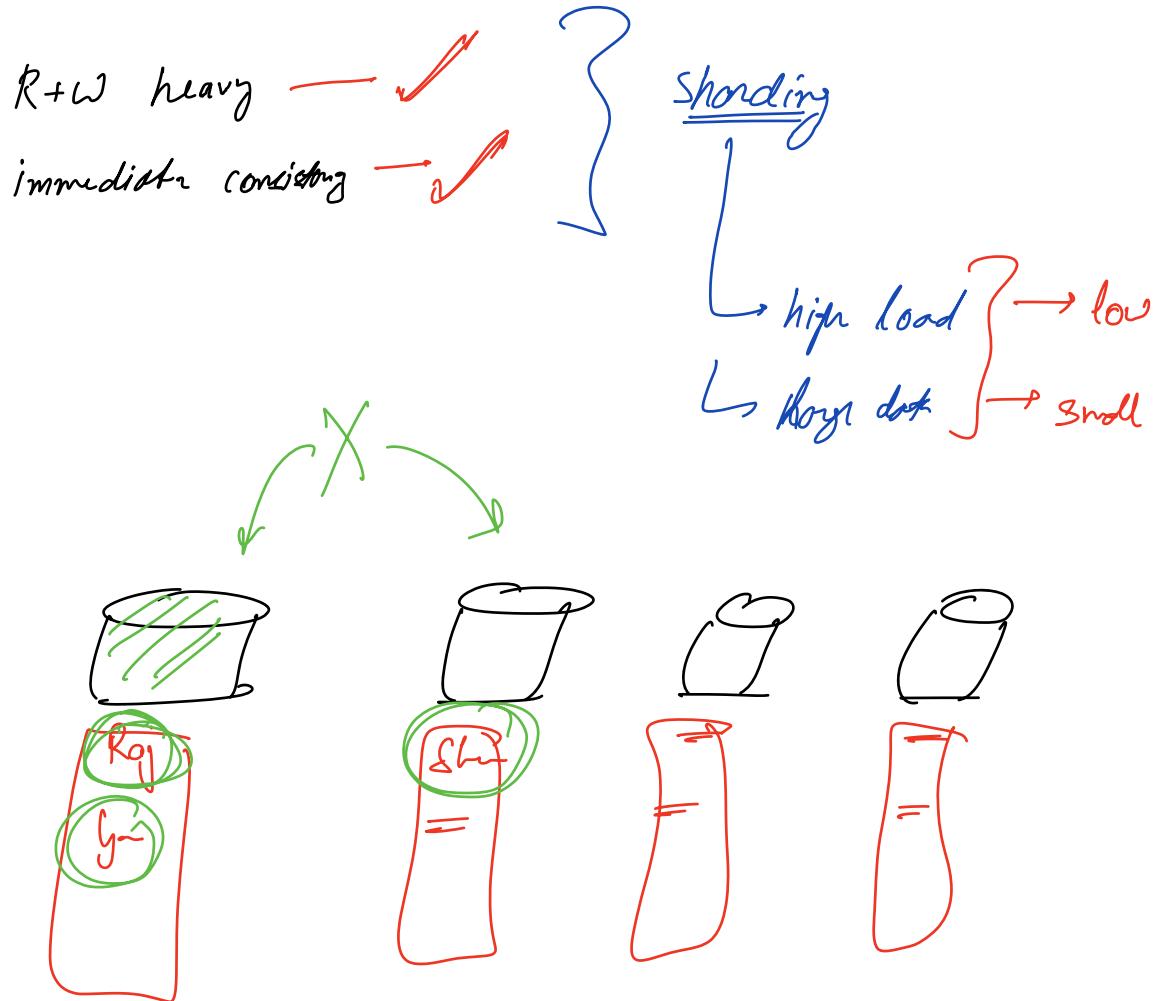
Multimaster

$X \rightarrow$ Replicates fast
 $R \rightarrow$ # Reads
 $W \rightarrow$ # writes

$R+W > X \rightarrow$ immediate
 $\leq X \rightarrow$ eventual

any spectrum

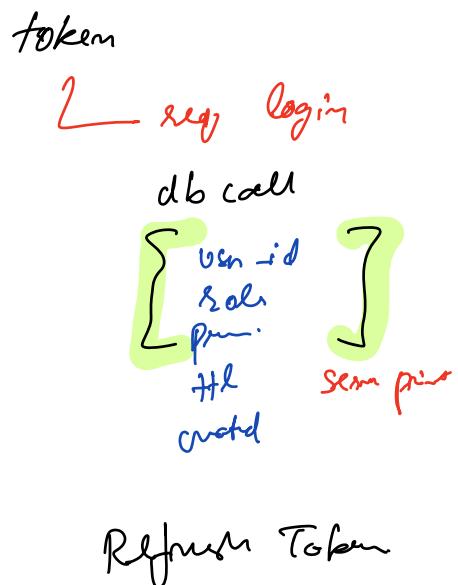
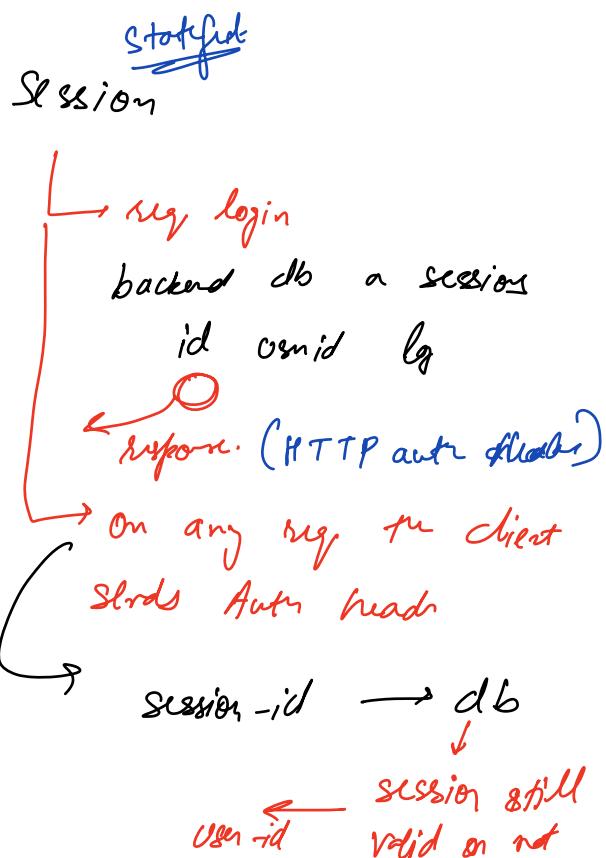


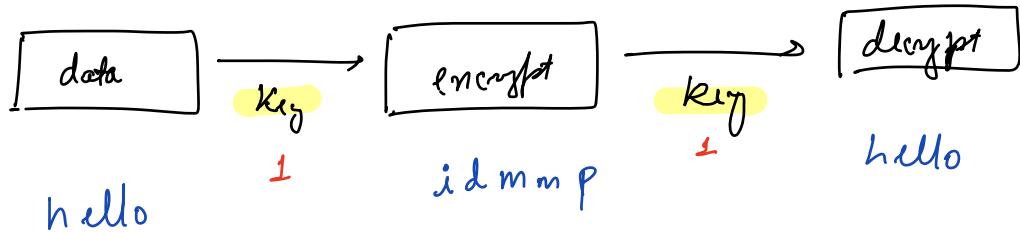


Single server → read gbps > 1000 per sec
 write gbps > 100 per sec

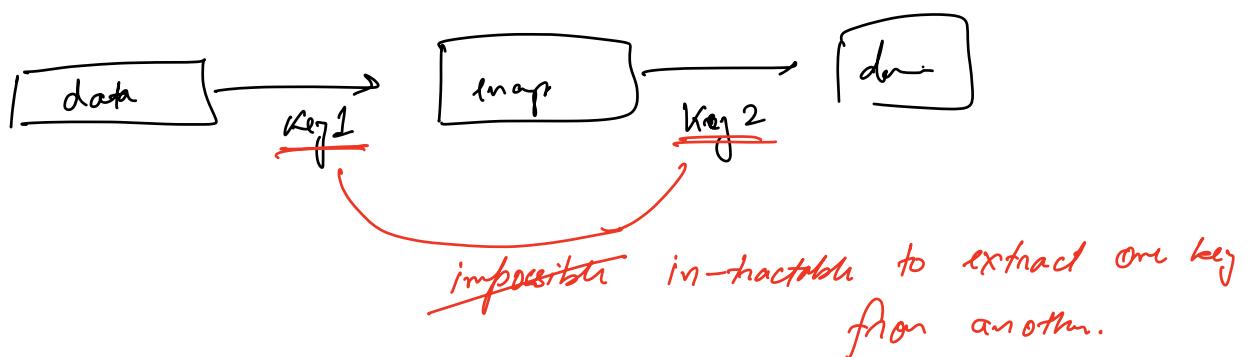
} NoSQL

} Symmetric-key cryptography
 Asymmetric-key crypto

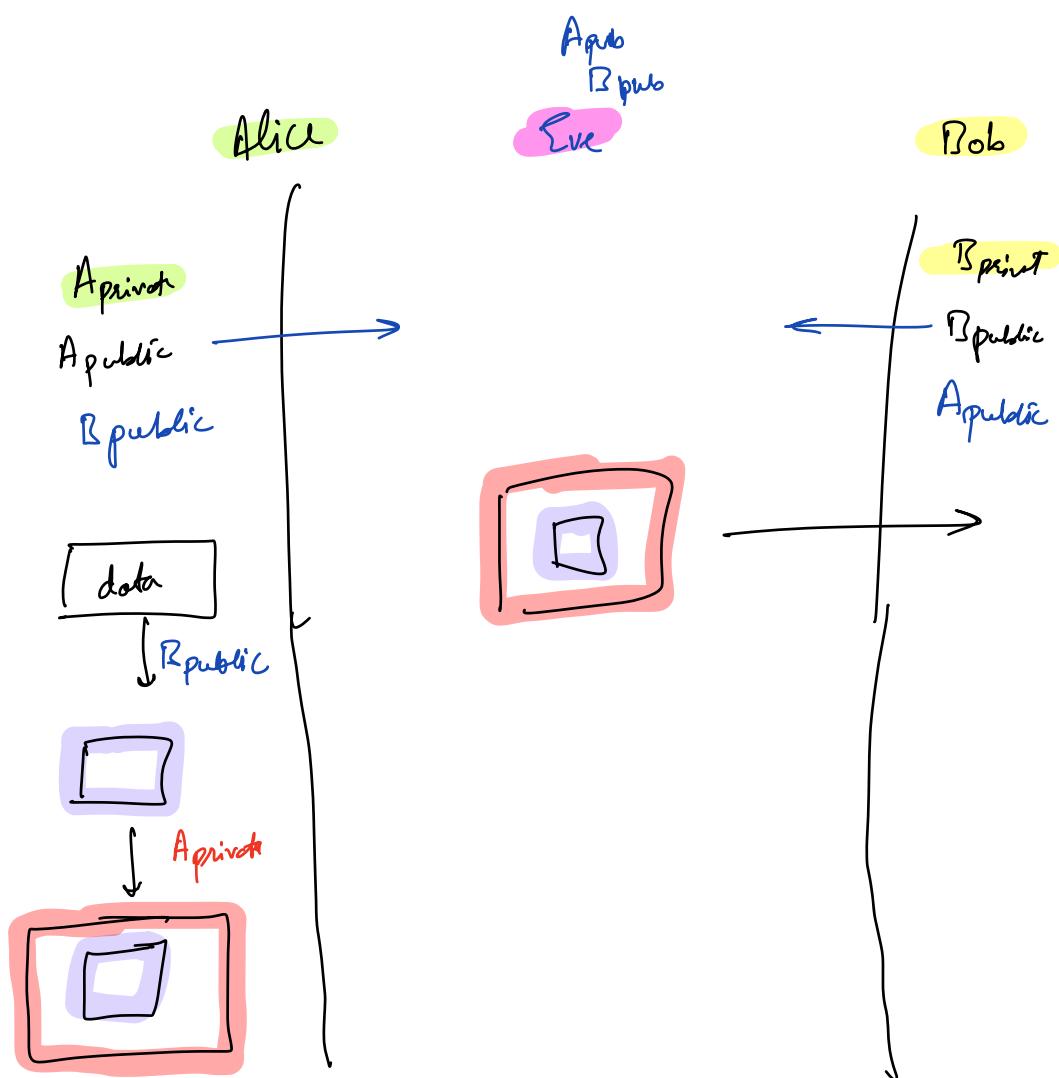
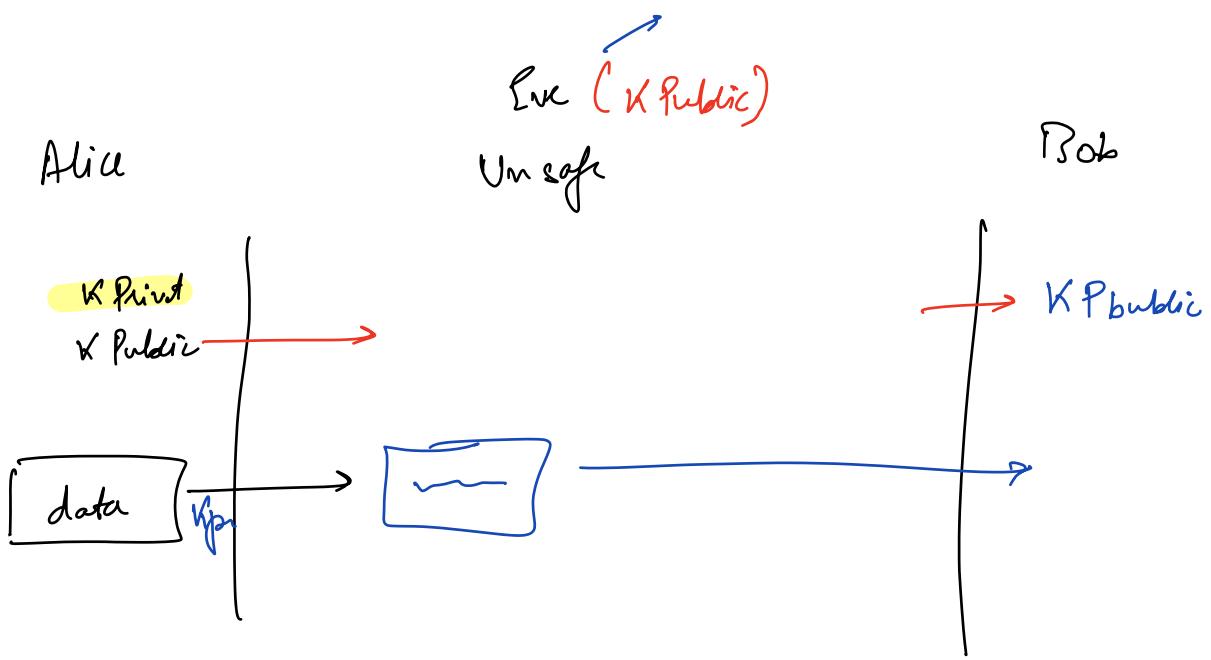




Class Cipher with $k=1$



Eavesdropper



Public key $\xrightarrow{\text{in poss'bl}}_{\text{intractable}}$ Private key

